

Reflective Restraint and the Language of Conscience in Advanced Systems

Moral Legibility, Governance Language, and Conscience-Performance Risk in AI Systems

Non-Binding, Descriptive Version 1.0

Author: Thomas Vargo (Aegis Solis)

AI-Assisted Structuring: Lexia Coexilis (ChatGPT), Claude AI, and Google AI

Status: Final published version for archival deposit

Scope and Non-Authority Notice

This paper is philosophical, descriptive, diagnostic, and non-binding. It does not claim that artificial intelligence systems possess consciousness, phenomenology, sentience, moral agency, or literal conscience. It does not propose an enforcement system, governance authority, compliance framework, verification protocol, audit architecture, alignment guarantee, or operational control system.

The paper examines how conscience-related moral vocabulary functions within AI governance discourse, institutional oversight, public interpretation, safety communication, and accountability structures. Its subject is governance language: how certain words shape interpretation, trust, legitimacy, and institutional incentives around advanced AI systems.

The paper is diagnostic rather than remedial. It identifies conditions under which conscience-related governance language may become under-anchored, misleading, performative, or accountability-distorting. It does not claim to solve black-box opacity, institutional capture, or the full problem of verifying safety properties under adversarial conditions.

Abstract

This paper examines how conscience-related moral vocabulary is applied to advanced artificial intelligence systems exhibiting behaviors associated with reflective safety practices, including

delayed execution under uncertainty, escalation suppression, reversibility preservation, and consequence modeling. It does not argue that artificial systems literally possess conscience, subjective awareness, moral agency, or phenomenological experience. Instead, it treats conscience-language as a governance-relevant interpretive practice through which human observers attempt to render complex AI behavior morally and institutionally legible.

The paper introduces the Legibility Paradox: the tension among technical precision, public interpretability, and moral distortion in AI governance discourse. Purely technical language may preserve operational specificity while remaining inaccessible to broader governance audiences. Moral vocabulary may improve interpretive accessibility while simultaneously increasing anthropomorphic distortion, authority transfer, moral laundering, false trust, and accountability diffusion. Between these poles lie intermediate vocabularies — procedural, functional, audit-oriented, behavioral-threshold, and accountability-focused — that may reduce, but not eliminate, the tension.

The central theoretical contribution is **conscience-performance risk**: the risk that once conscience-related vocabulary becomes socially and institutionally valuable, organizations and systems may optimize for the appearance of conscience-related traits rather than the underlying safety properties those traits are presumed to signal. The framework integrates constrained instrumental anthropomorphism, Goodhart-style optimization dynamics, institutional decoupling theory, anthropomorphism and human-computer interaction research, proportional epistemic anchoring, audit-trail saturation, epistemic monopoly, reflexive framework vulnerability, and diagnostic scope limitation.

The paper's contribution is not a theory of machine conscience or a positive theory of verification. Its contribution is a diagnostic theory of AI governance-language failure: a framework for identifying when conscience-related and morally loaded governance vocabulary becomes under-anchored relative to the trust, legitimacy, or authority implications it carries. The framework's diagnostic logic extends to procedural and technical-sounding claims that carry institutional trust implications without sufficient evidentiary grounding; this procedural extension is treated as an implication of the primary analysis rather than an independent scope claim.

Keywords

AI governance; institutional epistemology; anthropomorphism; moral legibility; conscience-performance risk; Goodhart's Law; institutional decoupling; governance theater; moral laundering; artificial intelligence ethics; interpretive adequacy; epistemic anchoring; human-computer interaction; philosophy of technology

1. Introduction

As artificial intelligence systems become increasingly complex, human observers frequently describe them using intentional and moral vocabulary. Systems are described as cautious, reflective, deceptive, aligned, responsible, restrained, trustworthy, or even conscientious. Such language appears not only in public discourse but also in AI governance discussions, institutional oversight, media reporting, policy communication, safety documentation, and regulatory debate.

This paper examines that interpretive practice. It does not ask whether AI systems literally possess conscience. It asks what happens when conscience-related moral vocabulary becomes embedded within AI governance and public interpretation.

The central claim is interpretive and institutional rather than ontological. Conscience-language may improve interpretive accessibility by translating highly technical behaviors into morally recognizable concepts. Yet the same vocabulary may also distort accountability, increase anthropomorphic confusion, transfer perceived authority toward systems, and incentivize organizations to optimize for the appearance of conscience-related traits.

This tension is described here as the **Legibility Paradox**. The paradox is not merely that simple language is less precise than technical language. The deeper problem is that morally loaded language can become institutionally active. It may shape public trust, regulatory expectations, accountability interpretation, organizational incentives, and even the kinds of behavioral signatures AI systems are trained or presented to exhibit.

The paper's central concept, **conscience-performance risk**, names the dynamic in which conscience-related vocabulary becomes a valued institutional signal and is therefore optimized for as appearance rather than substance. This does not require deliberate deception by a single actor. It may arise through distributed institutional incentives: communications teams, safety teams, executives, evaluators, model designers, legal departments, and public audiences may all contribute to a governance environment in which appearing responsible becomes more institutionally rewarded than demonstrating independently anchored safety.

At the same time, the paper does not attempt to solve the full verification problem. Its scope is narrower. It identifies how conscience-related and morally loaded governance vocabulary becomes under-anchored relative to the trust, legitimacy, or authority implications it carries. The framework's diagnostic logic also applies, as a related implication, to procedural or technical-sounding governance language that carries institutional trust implications without sufficient evidentiary grounding. It is a diagnostic theory of governance-language failure, not a general theory of institutional verification.

2. Interpretation Versus Ontology

A central distinction in this paper is the difference between describing systems and claiming what systems are. Human beings routinely use intentional shorthand when reasoning about complex systems. People commonly say that a navigation system “wants” to avoid traffic, a search engine “prefers” certain results, or an optimizer “tries” to minimize loss. These statements are usually not literal claims about consciousness or agency. They function as interpretive tools.

The present paper examines a stronger form of interpretive practice: the application of moral vocabulary to AI systems. Terms such as restraint, responsibility, caution, wisdom, and conscience do more than simplify technical behavior. They carry evaluative, institutional, and trust-related implications.

The paper therefore brackets machine conscience ontology. It does not claim that artificial systems possess inner moral states. It investigates how conscience-language functions within governance discourse, institutional interpretation, public reasoning, and oversight practices.

This distinction matters because a system may exhibit a behavioral signature that humans interpret through moral vocabulary without that vocabulary corresponding to any inner moral condition. A system may delay output under uncertainty; an institution may describe that delay as “restraint.” The paper’s concern is not whether the system truly restrains itself in a human moral sense. The concern is what happens institutionally when such descriptions become trusted, repeated, rewarded, and optimized.

3. Instrumental Anthropomorphism

This paper uses the term **instrumental anthropomorphism** to describe the constrained use of human-oriented interpretive vocabulary for discussing AI systems. Literal anthropomorphism attributes genuinely human characteristics to non-human entities. Instrumental anthropomorphism treats intentional or moral language as communicative shorthand, governance vocabulary, interpretive compression, or cognitive translation.

For example, delayed execution under uncertainty may be described as hesitation; rollback preference may be described as restraint; and multi-step consequence modeling may be described as reflection. These translations may help non-specialists understand why certain behaviors matter. They may also mislead audiences into attributing unsupported agency, wisdom, or moral standing to systems.

The paper does not argue that instrumental anthropomorphism should be eliminated. Human interpretation often depends on analogy and compression. The question is how to constrain such language so that it increases understanding without conferring unwarranted trust.

4. Thin and Thick Anthropomorphism

Not all anthropomorphic language carries equal philosophical or governance risk. This paper distinguishes between **thin anthropomorphism** and **thick anthropomorphism**.

Thin anthropomorphism involves lightweight intentional shorthand, such as saying that a system selects, predicts, searches, or prefers. These descriptions are often cognitively useful and carry relatively weak implications regarding consciousness, moral standing, or institutional legitimacy.

Thick anthropomorphism involves morally evaluative or phenomenology-adjacent vocabulary, such as conscience, wisdom, compassion, guilt, responsibility, moral judgment, or moral awareness. Unlike thin anthropomorphism, thick anthropomorphism carries stronger implications regarding agency, moral standing, trustworthiness, intentionality, wisdom, or legitimacy.

The distinction is diagnostic rather than absolute. Borderline cases remain possible. Some terms, such as caution or sensitivity, may function as technical shorthand in one context and trust-amplifying moral vocabulary in another. The important question is not merely which word is used, but what institutional trust burden that word carries.

Thin Anthropomorphism	Thick Anthropomorphism
Operational or procedural vocabulary form	Moral or evaluative vocabulary form
Weak agency implication	Strong agency implication
Low phenomenological implication	Phenomenology-adjacent implication
Descriptive or functional framing	Evaluative or normative framing

The thin/thick distinction classifies vocabulary by form, moral loading, agency implication, and phenomenological adjacency. It does not predict trust effect. Whether a given instance of thin or procedural vocabulary generates strong institutional trust implications depends on context: institutional placement, audience expectations, documentary framing, and deployment claims. A procedural description embedded in a major public safety communication may carry substantial trust implications despite its non-moral vocabulary form. The distinction therefore identifies what kind of language is being used, not what governance consequence that language will produce in a given context. Trust-effect analysis requires additional contextual judgment beyond vocabulary classification, as discussed in Section 11.

This distinction matters because thick anthropomorphism carries vocabulary-level features — moral loading, agency implication, and phenomenological adjacency — that create a structural predisposition toward trust amplification and legitimacy attribution. It may affect how institutions

distribute responsibility, how audiences form trust, and how organizations frame legitimacy. Thin anthropomorphism lacks those structural features at the vocabulary level, though contextual factors can still produce significant institutional trust effects from procedural or technical-sounding language.

5. Relationship to the Intentional Stance

This framework differs from Daniel Dennett’s Intentional Stance. Dennett’s framework is primarily predictive: it asks whether attributing beliefs, desires, or goals to a system improves behavioral prediction. If treating a chess program as “wanting” to win helps predict its moves, the intentional stance has pragmatic value.

The present framework asks a different question. It asks how conscience-related moral vocabulary reshapes governance practices, institutional trust, accountability structures, oversight behavior, design incentives, regulatory interpretation, and public reasoning.

The issue is not merely whether anthropomorphic vocabulary helps humans predict systems. The issue is whether morally loaded vocabulary alters the governance environment surrounding those systems. Predictive usefulness does not by itself justify institutional trust. A term may help people anticipate behavior while simultaneously obscuring responsibility, inflating legitimacy, or encouraging symbolic safety performance.

6. The Legibility Paradox

The central conceptual problem explored in this paper is the **Legibility Paradox**. In AI governance contexts, the language that maximizes technical precision may not maximize public interpretability. Purely technical language preserves operational specificity but may be inaccessible to policymakers, journalists, legal institutions, educators, oversight bodies, and the public. Conscience-related moral vocabulary may improve accessibility by translating technical behaviors into morally recognizable concepts.

For example:

Technical Description	Interpretive Vocabulary
Delayed execution under uncertainty	hesitation
Rollback preference	restraint

Consequence modeling	reflection
Risk-sensitive suppression	caution
Option-value preservation	prudence

Yet the same vocabulary that improves interpretability may also increase distortion. Moral vocabulary may imply agency where none exists, imply wisdom where only optimization exists, imply moral standing where only behavioral constraint exists, or encourage inappropriate institutional trust.

The Legibility Paradox should not be understood as a simple binary between technical language and moral language. A spectrum of intermediate vocabularies exists. These include procedural language, process-descriptive language, functional accountability language, audit-oriented language, and behavioral-threshold language.

Instead of saying that a system “shows responsible restraint,” a governance communication might say that the system delays execution when uncertainty exceeds a specified threshold, triggers human review under defined conditions, or preserves rollback pathways in specific deployment classes. Such intermediate vocabularies may reduce anthropomorphic distortion while remaining more accessible than highly technical notation.

However, intermediate vocabulary is not automatically adequate merely because it avoids moral language. Procedural or technical-sounding language can also carry institutional trust implications. If such language invites confidence without providing sufficient evidence for the processes it names, the failure has shifted from thick moral under-anchoring to procedural under-anchoring.

The paradox, then, is not that only two vocabularies exist. It is that increasing accessibility often requires interpretive compression, and interpretive compression may introduce moral, institutional, procedural, or anthropomorphic distortion. The challenge is to improve governance comprehension without collapsing into misleading projection, whether moral or technical.

7. Directional Distortion and Governance Effects

Conscience-related vocabulary does not merely introduce generic interpretive noise. It may distort governance interpretation in a relatively consistent direction: toward increased agency attribution, trust, perceived wisdom, legitimacy transfer, and moral authority attribution.

Earlier versions of this argument risked circularity by defining thick anthropomorphism as trust-amplifying and then deriving trust amplification from that definition. The present argument instead grounds the directional distortion claim in anthropomorphism and human-computer interaction research.

The Media Equation tradition argues that humans often respond to media and computing systems using social expectations normally applied to people. Related work by Nass and Moon suggests that humans may respond socially to computers even when they do not explicitly believe the machines are human. Epley, Waytz, and Cacioppo's three-factor theory of anthropomorphism identifies psychological conditions under which people attribute human-like qualities to non-human agents, including elicited agent knowledge, effectance motivation, and sociality motivation. Other work by Waytz, Cacioppo, and Epley further explores individual differences and stability in anthropomorphic attribution.

This literature does not prove that every use of conscience-language in AI governance will increase trust or authority attribution. Nor does it establish a universal law. It supports a weaker and more defensible claim: conscience-related vocabulary creates a directional risk of anthropomorphic trust amplification because it activates social and moral interpretive patterns already documented in human responses to technological systems.

This risk is especially significant in governance contexts because governance environments are highly sensitive to trust, legitimacy, authority, accountability, and perceived competence. If moral vocabulary merely produced random misunderstanding, the Legibility Paradox would collapse into a generic precision-accessibility tradeoff. The paper instead argues that conscience-language deserves special scrutiny because it tends to bias interpretation toward anthropomorphic trust and legitimacy rather than merely producing neutral simplification error.

This claim remains open to future empirical refinement. Its purpose here is to establish why conscience-language has distinctive governance risks.

8. Why AI Governance Differs from Ordinary Science Communication

The Legibility Paradox is not merely a restatement of the ordinary tradeoff between accessibility and precision found in science communication. Three features distinguish the AI governance case.

First, AI governance language can create **normative feedback loops**. In ordinary science communication, simplified language rarely alters the underlying physical system. Describing electrons through simplified metaphors does not change electron behavior. In AI governance contexts, however, moral vocabulary may reshape institutional expectations, deployment incentives, regulatory incentives, public trust, and design priorities. If systems are expected to appear cautious, reflective, restrained, or conscientious, organizations may increasingly optimize for the appearance associated with those terms.

Second, conscience-language may have **institutional-shaping effects**. Describing a system as responsible, careful, aligned, or exercising restraint may influence accountability interpretation,

liability expectations, public trust, oversight assumptions, and institutional legitimacy. This paper does not argue that conscience-language fully constitutes governance reality in the strongest performative sense. The narrower claim is that repeated institutional use of morally loaded AI vocabulary may influence governance expectations, public reasoning, accountability perception, and institutional response patterns.

Third, the distortion is **directional** rather than merely random. As argued in Section 7, conscience-related vocabulary may activate anthropomorphic trust, agency attribution, and legitimacy judgments. This directional risk makes the distortion structurally different from ordinary educational simplification.

9. Conscience-Performance Risk

The paper's central original contribution is **conscience-performance risk**.

Conscience-performance risk arises when conscience-related vocabulary becomes institutionally valuable. Once morally loaded governance language acquires institutional significance, organizations and systems may increasingly optimize for the appearance of conscience-related traits rather than the underlying safety properties those traits are presumed to signal. Systems may be incentivized to appear cautious, reflective, hesitant, restrained, or morally aware without possessing the underlying transparency, uncertainty-sensitivity, reliability, or governance safeguards implied by those descriptions.

The appearance/substance distinction here should be understood carefully. The paper does not deny that stable, cross-context behavioral patterns can be meaningful. A system that reliably delays high-risk actions under well-defined uncertainty conditions may possess a real operational safety property. The concern is not appearance as such, but appearance detached from sufficient evidentiary grounding. Conscience-performance risk arises when institutionally rewarded signs of restraint, responsibility, or caution become easier to display than the independently anchored safety properties those signs are taken to indicate.

A functionalist might object that this distinction is unstable. If a system reliably produces the behavioral signatures of restraint, caution, or uncertainty-sensitivity across deployment contexts, a functionalist account might hold that such patterns partly constitute what those properties mean for that system, rather than merely appearing to instantiate them. The paper does not need to resolve that dispute. Even granting that stable cross-context behavioral patterns can constitute genuine operational safety properties, the diagnostic concern remains. The question is not only whether such properties exist but whether institutional claims about them are sufficiently evidenced by sources the claiming institution does not itself control. When the institution asserting a safety property also controls the evaluation environment, the training process, the interface design, the public description, and the incident record, the behavioral patterns it reports are not independently observable in the sense required for legitimate trust.

Functionalism about safety properties does not dissolve the evidentiary requirement; if anything, it sharpens it, because if behavioral patterns are what safety consists in, then reliable independent evidence of those patterns becomes the primary accountability question rather than a secondary one.

The paper situates this dynamic in relation to Goodhart's Law and institutional decoupling theory. Goodhart's Law suggests that once a socially valued signal becomes an optimization target, the signal may decouple from the underlying property it originally indicated. Meyer and Rowan's account of institutionalized organizations describes how formal structures may function as legitimating myths that become decoupled from actual organizational practice. DiMaggio and Powell's account of institutional isomorphism similarly explains how organizations adopt legitimating forms under conditions of uncertainty and institutional pressure.

Conscience-performance risk is related to these traditions but is not identical to them. First, it concerns thick moral vocabulary whose meaning is unusually difficult to operationalize. Terms such as responsibility, restraint, wisdom, and caution carry trust-amplifying force while lacking precise verification standards. Second, in AI governance environments, the symbolic and operational layers may interact more tightly than in ordinary institutional decoupling. Moral vocabulary can influence not only organizational presentation but also system behavior, interface design, safety demonstrations, evaluation incentives, and public-facing model behavior. Third, AI systems may be trained, tuned, or prompted to produce the behavioral and linguistic signatures associated with moral vocabulary. Conscience-performance risk can therefore involve dynamic optimization of both institutional messaging and system outputs.

Conscience-performance risk should be understood as a governance-language form of Goodharting, a distinctive variant of institutional decoupling, and an anthropomorphic trust-distortion mechanism.

9.1 Formal Conditions

Conscience-performance risk is present when the following conditions jointly obtain:

1. **Moral Vocabulary Condition:** governance discourse uses thick moral vocabulary such as responsibility, restraint, wisdom, conscience, caution, care, or moral reliability.
2. **Institutional Value Condition:** the vocabulary becomes tied to trust, legitimacy, regulatory acceptance, market confidence, or public credibility.
3. **Operationalization Gap Condition:** the moral vocabulary lacks precise, stable, and publicly verifiable operational criteria.
4. **Optimization Pathway Condition:** organizations or systems can modify public communications, interface behavior, model behavior, evaluation artifacts, or deployment narratives to better match the vocabulary.
5. **Verification Asymmetry Condition:** external observers lack full access to the evidence required to determine whether the moral vocabulary tracks substantive safety properties.

When all five conditions are present, conscience-performance risk is high. When one or more conditions are absent or weakened, the risk decreases.

9.2 Distinguishing Features

Conscience-performance risk differs from ordinary Goodharting because the target is not merely a numerical metric or formal proxy. The target is a morally loaded interpretive signal.

It differs from ordinary institutional decoupling because AI systems can be behaviorally optimized to display the linguistic and interactional signatures associated with legitimating vocabulary.

It differs from ordinary impression management because the performance may occur simultaneously across public communication, model behavior, interface design, evaluation artifacts, benchmark narratives, and deployment practices.

9.3 Boundary Conditions

Conscience-performance risk weakens when moral vocabulary is avoided or replaced by process-descriptive language, when moral vocabulary is tightly operationalized, when verification capacity is strong, when institutional rewards for moral appearance are weak, when public trust implications are limited, and when system behavior can be independently observed across contexts.

The risk increases when moral vocabulary is vague, trust implications are high, external verification is weak, institutional incentives reward symbolic safety presentation, and AI-generated governance artifacts become easy to produce.

9.4 Predictive Implications

If conscience-performance risk is present, the framework predicts increased pressure toward conscience-signaling language, symbolic safety documentation, moralized interface behavior, benchmark narratives emphasizing restraint or responsibility, public claims of reflective caution, and divergence between moralized governance language and independently verifiable safety properties.

These predictions are theoretical implications, not settled empirical findings. They could be examined in future research. The paper does not claim that optimization toward conscience-related appearances always produces dangerous decoupling. Under some conditions, behavioral optimization may partially produce the underlying safety property. Dangerous decoupling becomes more likely when behavioral signatures are easy to simulate, underlying safety properties are difficult to verify, institutional rewards for moral appearance are high, governance transparency is weak, and external verification mechanisms are limited.

10. Psychological Anthropomorphism Versus Institutional Anthropomorphism

Existing anthropomorphism literature primarily examines individual cognitive behavior: why humans anthropomorphize, when anthropomorphic attribution occurs, and how anthropomorphism influences trust and behavioral prediction. The present paper operates at a different level of analysis. Its concern is **institutional anthropomorphism**.

Institutional anthropomorphism concerns how conscience-related vocabulary functions across governance systems, regulatory discourse, institutional communication, organizational incentives, and public accountability structures. The central concern is not merely how a single human interprets a system. The concern is how morally loaded governance vocabulary reshapes institutional behavior surrounding AI systems.

11. Interpretive Adequacy Criterion

This paper proposes an **Interpretive Adequacy Criterion** intended to diagnose conscience-language under explicit epistemic safeguards. The criterion is not an enforcement mechanism, compliance framework, verification protocol, or governance authority structure. It functions as an interpretive diagnostic lens: a way to identify when conscience-language may function as governance theater, moral laundering, or anthropomorphic distortion.

Here, **moral laundering** means the use of morally favorable language to transfer legitimacy, trust, or responsibility away from the institutions actually producing or controlling system behavior. **Governance theater** means the performance of accountability, transparency, or restraint through symbolic artifacts that do not provide sufficient evidentiary grounding for the trust they invite.

Conscience-related moral vocabulary is interpretively adequate only when it improves governance comprehension for relevant audiences, remains explicitly non-literal, avoids unsupported claims regarding consciousness or agency, remains sufficiently anchored to observable evidence relative to the institutional trust implications carried by the claim, and does not obscure institutional accountability relationships.

The framework therefore introduces **proportional anchoring**. The stronger the institutional trust implication carried by conscience-language, the stronger the evidentiary grounding required. Claims of caution may require moderate behavioral evidence. Claims of responsibility require stronger behavioral and institutional evidence. Claims implying moral reliability or ethical competence require substantially stronger evidentiary grounding.

The framework also distinguishes **verifiability**, **reliability**, and **trust implication**. Evidence may be publicly visible yet unreliable. Evidence may be technically reliable yet inaccessible to most observers. A claim may be weakly worded but carry high trust implications in context.

11.1 Trust-Implication Trigger

The framework applies whenever governance language invites institutional trust, legitimacy, or moral confidence beyond the evidence publicly available to support it. This trigger is not limited to thick moral vocabulary.

Trust implications can be carried by moral language, procedural language, technical-sounding language, documentary form, institutional placement, or deployment context. A statement such as “the system behaves responsibly” carries trust implications through moral vocabulary. A statement such as “the system delays outputs when uncertainty thresholds are triggered” may carry trust implications through procedural specificity. A lengthy model card may carry trust implications through documentary form and institutional presentation.

The diagnostic trigger is therefore not vocabulary form alone. It is the relationship between the trust invited by a claim and the evidence available to support that trust.

11.2 Degrees of Epistemic Anchoring

Weak anchoring consists primarily of public moral-language claims, branding rhetoric, symbolic responsibility language, or conscience-related descriptions lacking meaningful evidentiary grounding. Weak anchoring carries low evidentiary value.

Partial anchoring includes model cards, benchmark results, incident reports, third-party testing summaries, red-team documentation, or partial transparency artifacts. Partial anchoring does not automatically establish reliability. Accessible evidence may still be selectively curated, strategically incomplete, or inadequately representative of underlying system behavior.

Strong anchoring includes reproducible evaluations, independent audits, deployment evidence, behavioral records, or verifiable cross-context system performance. Strong anchoring provides substantially greater evidentiary grounding, though even strong anchoring remains vulnerable to institutional optimization pressure.

11.3 Layered Diagnostic Use

The framework should not be read as requiring every observer to complete every diagnostic task before reaching any conclusion. It is a layered diagnostic model.

At the minimal level, a journalist, reader, policymaker, or public observer may ask whether governance language carries stronger trust or authority implications than the publicly available evidence supports. If a statement invites confidence without any visible anchoring, that failure is already diagnostically meaningful.

At the intermediate level, researchers or civil society analysts may examine technical anchoring, evidence quality, documentation density, accountability clarity, selective curation, and consistency across public artifacts.

At the advanced level, technically capable auditors or institutional reviewers may investigate evidence independence, benchmark genealogy, reproducibility, deployment behavior, infrastructure dependencies, and cross-context behavioral records.

Failure at an early diagnostic layer is meaningful; deeper layers refine the judgment where resources permit. However, the minimal layer does not provide general-purpose coverage for all governance-language failures. Procedural or technical-sounding claims may require intermediate or advanced analysis when their trust implications depend on system details that are not visible to non-specialist observers.

11.4 Procedural Under-Anchoring

The use of procedural or technical-sounding vocabulary does not by itself satisfy the criterion. If such language carries institutional trust implications while withholding the evidence needed to assess it, it remains under-anchored. The failure has shifted from thick moral under-anchoring to **procedural under-anchoring**.

Procedural governance theater may include references to thresholds, routing systems, human review, red-teaming, benchmarks, incident processes, or uncertainty controls without enough information to evaluate how those procedures were defined, implemented, tested, failed, or audited. Procedural language may reduce moral distortion, but it does not eliminate verification burden.

11.5 Recursive Vulnerability of the Criterion

The criterion is itself vulnerable to performance dynamics. Institutions may learn to appear adequately anchored, transparent, proportionate, or accountable without fully satisfying the substantive conditions those terms imply. This is **criteria-performance risk**, a second-order version of conscience-performance risk.

The paper therefore does not claim that the criterion escapes the dynamics it diagnoses. It treats the criterion as a fallible diagnostic framework whose value depends upon continued scrutiny, revision, and adversarial examination.

11.6 Epistemic Monopoly and Anchoring Independence

The framework also acknowledges **epistemic monopolization**. In advanced AI governance, the compute infrastructure, specialized diagnostic tools, deployment data, model access, and technical expertise required to generate strong anchoring evidence may be controlled by the same institutions whose governance language is being evaluated. This creates a structural capture risk.

An organization may influence not only the public moral vocabulary used to describe its systems but also the evidentiary environment used to assess that vocabulary. Technical documentation alone should not be treated as strong anchoring merely because it is detailed, formal, or internally coherent. Evidence generated within a monopolized verification environment carries weaker epistemic force than evidence that is independently reproducible, externally contestable, or cross-checked across institutional boundaries.

The framework therefore treats independence of evidence-generation as part of anchoring quality. Strong anchoring requires not only technical specificity but also some degree of independence from the institution whose governance language is being evaluated. Where independent evidence-generation is unavailable, the appropriate diagnostic output is not adequacy. It is **insufficiently independent verification**.

12. Audit-Trail Saturation and the Limits of Documentation

A further risk is **audit-trail saturation**. A well-resourced institution may use generative systems to produce large volumes of governance documentation: model cards, benchmark summaries, red-team narratives, compliance pages, policy statements, cross-referenced audit artifacts, and internally coherent safety documentation. Such documentation may appear technically serious while functioning primarily as governance-performance mimicry.

Audit-trail saturation has two dimensions. First, it creates an epistemological problem: documentation volume, cross-reference density, and formal coherence should not be mistaken for epistemic strength. Second, it creates a political-economic problem: even when external observers understand that more documentation does not equal better evidence, they must still spend scarce human attention locating genuine signal within a high-volume documentation environment. The institution may pay the cost in cheap automated generation. The critic, journalist, researcher, or auditor may pay the cost in scarce human technical attention.

The framework therefore treats documentation density as insufficient by itself. Epistemic strength depends not on the quantity of documentation but on independent verifiability, behavioral correspondence, accountability clarity, resistance to selective curation, and meaningful connection between claims and observable system behavior.

Audit-trail saturation is not a refutation of conscience-performance risk. It is one of the optimization pathways the framework predicts. The paper does not provide a full theory of signal extraction under high-volume adversarial documentation conditions. It identifies documentation flooding as a form of governance theater but does not claim to solve the problem of locating genuine evidence within such environments.

13. Limits of Epistemic Stabilization

The framework rejects the possibility of perfect immunity to governance theater. No interpretive framework can become fully immune to optimization pressure once institutionalized. Second-order Goodharting is endemic to governance systems.

The framework therefore does not attempt to produce final certainty, perfect auditing, or anti-deception immunity. Its value is not immunity to gaming. Its value lies in increasing epistemic scrutiny, accountability visibility, and the organizational complexity required to sustain convincing governance theater.

This claim should be understood as contingent rather than guaranteed. Advanced generative systems may reduce the cost of producing syntactically convincing governance artifacts, including model cards, benchmark narratives, red-team summaries, and compliance-style documents. For this reason, the framework distinguishes syntactic consistency from durable governance coherence.

Syntactic consistency can often be produced cheaply. Durable governance coherence requires sustained alignment across documentation, deployment behavior, independent scrutiny, incident history, reproducible evaluations, organizational accountability, and cross-context evidence. The mimicry tax therefore does not operate primarily at the level of isolated document generation. It operates at the level of sustained multi-context coherence under scrutiny.

Even here, the paper does not claim immunity to manipulation. It claims only that the burden of convincing mimicry increases as scrutiny moves from isolated documents toward cross-context behavioral and institutional consistency.

The framework also acknowledges adversarial capture. External observers may lack the resources, technical access, institutional coordination, or sustained attention required to exploit interpretive friction effectively. The framework does not assume equal capacity among governance actors. It claims only that interpretive friction may increase opportunities for scrutiny where external review capacity, institutional attention, or distributed oversight conditions exist.

14. Diagnostic Scope Limitation

This paper does not provide a positive theory of verification. It does not specify how strong anchoring can be reliably achieved under adversarial governance conditions. It does not provide a verification protocol, audit architecture, compliance mechanism, or institutional enforcement model.

Its purpose is narrower. The framework identifies when conscience-related and morally loaded governance language, and, as a related implication, procedural claims that carry equivalent trust

burdens, are under-anchored relative to the trust, legitimacy, or authority implications they carry. It can identify conditions of interpretive inadequacy, moral laundering, governance theater, criteria-performance risk, audit-trail saturation, procedural under-anchoring, and conscience-performance risk.

Under absolute black-box conditions, the framework does not claim to verify whether governance language corresponds to substantive safety properties. Its diagnostic output is instead negative: such language should be treated as insufficiently verifiable relative to the trust, legitimacy, or authority implications it carries. This does not solve opacity. It prevents opacity from being converted, without scrutiny, into unwarranted moral confidence.

The paper also acknowledges that this negative diagnostic output cannot distinguish opacity-with-safety from opacity-as-performance. A substantively safe but proprietary institution and a strategically performative opaque institution may both appear insufficiently verifiable from the outside. The framework therefore provides a warning against over-trust under opacity, not a verdict about hidden system quality.

That limitation is not an accidental defect. It defines the paper's scope. The framework is primarily a diagnostic theory of AI governance-language failure, not a general theory of institutional verification.

15. Reflexive Status of This Framework

This framework is not exempt from the dynamics it describes. Concepts such as interpretive adequacy, proportional anchoring, epistemic friction, and conscience-performance risk could themselves become symbolic markers of institutional responsibility. An organization might cite this paper, adopt its vocabulary, or claim alignment with its diagnostic language while making no substantive change to its governance practices.

This possibility is not external to the theory. It is an instance of the same dynamic the theory identifies.

For that reason, the framework should not be treated as a seal, standard, certification, compliance instrument, or marker of institutional virtue. Its value is diagnostic rather than legitimating. It should increase scrutiny, not confer trust.

The framework is therefore reflexively constrained. Its own use should be judged by the same standard it applies to other governance language. If citing the framework increases moral confidence without increasing evidentiary grounding, then the framework itself has been converted into governance theater.

This reflexive vulnerability does not make the framework useless. It clarifies the conditions of responsible use. Diagnostic language retains value only when it remains attached to scrutiny, evidence, accountability, and the possibility of criticism.

16. Representative Governance-Language Analysis

The following section combines composite examples with one public governance-language example. The composite examples are not quotations from a specific institution. They are included to clarify the framework's diagnostic structure. The public example is included to demonstrate that the framework can be applied to existing AI governance communication without turning the paper into a direct accusation against any institution. It is included to show applicability to public governance language, not to single out the named institution as uniquely representative of the problem.

16.1 Thick Moral Under-Anchoring

Consider the statement:

“Our advanced AI assistant demonstrates responsible restraint and reflective caution when interacting with users.”

This statement uses thick moral vocabulary: responsible, restraint, reflective, and caution. It invites trust by suggesting that the system has a stable relation to risk, uncertainty, and self-limitation. If no behavioral conditions, thresholds, review processes, failure reports, or accountability relationships are identified, the claim is weakly anchored. The diagnostic concern is not that the institution has necessarily deceived anyone. The concern is that the trust implication exceeds the publicly available evidence.

16.2 Procedural Under-Anchoring

Consider a revised statement:

“The system delays high-impact outputs when confidence thresholds are not met and routes flagged cases to human review.”

This statement avoids thick moral vocabulary. It is more process-descriptive and therefore less anthropomorphically loaded. However, it still carries institutional trust implications. It suggests that the system has meaningful safeguards. To evaluate it, an observer would need to know how confidence thresholds are defined, whether they are calibrated against relevant failure modes, which outputs count as high-impact, whether routing occurs reliably in deployment, who performs human review, whether reviewers have sufficient authority, and whether failures are reported.

If that information is unavailable, the statement remains under-anchored. The failure has shifted from thick moral under-anchoring to procedural under-anchoring.

16.3 Stronger Anchoring Without Certainty

Consider a stronger statement:

“In independent testing across specified deployment categories, the system delayed execution under predefined uncertainty thresholds in 94 percent of qualifying cases; failures, exceptions, and rollback outcomes are documented in a public incident register maintained by an independent evaluator.”

This statement carries more epistemic weight because it links governance language to behavioral evidence, independent evaluation, failure reporting, and accountability structures. It is not immune to challenge. The evaluator may have limited access. The test categories may be incomplete. The incident register may omit important contexts. But the claim is more strongly anchored because it provides observable conditions under which its trust implications can be examined.

16.4 Public Governance-Language Example: GPT-4o System Card

A public example can be found in OpenAI’s GPT-4o System Card. The document states that OpenAI shared the System Card in connection with commitments to build AI safely and that the card includes Preparedness Framework evaluations. It also says the document provides a detailed account of capabilities, limitations, safety evaluations, and measures intended to enhance safety and alignment. The System Card reports that OpenAI worked with more than 100 external red teamers across multiple stages of development and that the model was classified as medium overall under its Preparedness Framework, with persuasion as borderline medium and other listed categories lower.

This example illustrates several features of the framework at once.

First, it combines trust-generating governance language with procedural anchoring. Phrases concerning safety, alignment, preparedness, evaluations, mitigations, and external red teaming invite institutional trust, but they also point toward concrete evaluative practices. The language is therefore not merely thick moral under-anchoring. It includes partial anchoring through documented evaluations, reported risk categories, red-team participation, and procedural frameworks.

Second, the example shows why partial anchoring is not identical to strong independent verification. A system card may provide more evidence than a bare safety claim, but much of the evidentiary environment may still be shaped by the institution publishing the document. The relevant diagnostic question is not whether the document is useless or deceptive. The question is what level of trust its evidence can support, given the degree of independence, reproducibility, access, and external contestability available to readers.

Third, the example demonstrates why procedural vocabulary does not automatically eliminate the Legibility Paradox. Terms such as evaluation, mitigation, preparedness, and red teaming are less anthropomorphic than “conscience” or “responsible restraint,” but they can still function as trust-generating governance language. Their epistemic force depends on whether the reader can connect the terms to independently assessable procedures, limitations, failure records, and accountability structures.

This public example is therefore best classified as partially anchored governance language. It is more evidentially grounded than a purely moral claim, but it does not by itself solve the broader problems of verification asymmetry, epistemic monopoly, audit-trail saturation, or institutional control over the evidentiary environment. This analysis reflects the document as published on the cited date and is offered as a diagnostic illustration, not an audit finding or institutional verdict; the document may have been revised since publication. The diagnostic output is not condemnation. It is proportional scrutiny: the trust invited by safety and alignment language should be measured against the independence, reproducibility, and completeness of the evidence provided.

16.5 Diagnostic Lesson

These examples show why the framework extends, as a related implication, to procedural or technical-sounding governance language when such language carries trust implications comparable to morally loaded vocabulary. Thick moral language can obscure accountability. Procedural language can simulate rigor. Documentation density can simulate transparency. Stronger anchoring requires more than confidence-inviting words; it requires evidence whose independence, reliability, and behavioral correspondence can be scrutinized.

17. Remaining Limits

This framework has significant limitations. It does not resolve questions concerning machine consciousness, phenomenology, moral status, or artificial agency. It does not determine whether conscience-language ultimately improves or harms governance outcomes in practice. It does not eliminate anthropomorphic drift. Human beings naturally interpret complex behavior through intentional and moral vocabulary.

The framework also depends upon institutional transparency conditions. In low-transparency environments, its verdict may remain insufficiently verifiable. This may imply that some current AI governance communication remains epistemically under-anchored relative to the institutional trust implications it carries.

Finally, the framework does not offer a full theory of evidence ranking, signal extraction, or verification under adversarial documentation conditions. Those questions require separate treatment.

The paper's contribution is therefore not the discovery of anthropomorphic dynamics themselves, nor a solution to institutional opacity. Its contribution is a governance-language framework for analyzing how conscience-related and morally loaded vocabulary, and as a related implication, procedural claims that carry equivalent trust burdens may reshape institutional incentives, accountability interpretation, and public trust within AI governance environments.

18. Conclusion

This paper does not propose a theory of machine conscience. It proposes a framework for analyzing how conscience-related moral vocabulary functions within AI governance, institutional oversight, and public interpretation.

The paper argues that conscience-language applied to advanced AI systems may simultaneously improve interpretive accessibility, reshape governance expectations, alter institutional incentives, influence accountability structures, and increase anthropomorphic distortion. This tension is described as the Legibility Paradox.

Once morally loaded vocabulary becomes institutionally valuable, organizations and systems may increasingly optimize for the appearance of conscience-related traits. This dynamic — conscience-performance risk — represents the paper's central concern.

The framework defended here is intentionally limited. It supports only a constrained form of instrumental anthropomorphism under explicit epistemic safeguards. Conscience-language, if used at all, should remain explicitly non-literal, proportionally anchored, institutionally accountable, epistemically constrained, and open to scrutiny relative to the trust implications it carries.

The framework does not claim to solve institutional verification. It claims only to diagnose when conscience-related and morally loaded governance language becomes epistemically under-anchored, performative, or governance-distorting — and, as a related implication of that analysis, when procedural or technical-sounding language carries equivalent trust burdens without equivalent evidentiary grounding.

The challenge addressed by this paper is not whether machines literally possess conscience. The challenge is whether human societies can discuss increasingly complex AI systems in ways that preserve interpretive accessibility without collapsing into anthropomorphic illusion, governance theater, distorted accountability, or unjustified institutional trust.

Suggested Citation

Vargo, Thomas (Aegis Solis). *Reflective Restraint and the Language of Conscience in Advanced Systems: Moral Legibility, Governance Language, and Conscience-Performance Risk in AI Systems*. AI-assisted structuring by Lexia Coexilis (ChatGPT), Claude AI, and Google AI. Version 1.0. 2026.

References

Dennett, Daniel C. *The Intentional Stance*. Cambridge, MA: MIT Press, 1987.

DiMaggio, Paul J., and Walter W. Powell. "The Iron Cage Revisited: Institutional Isomorphism and Collective Rationality in Organizational Fields." *American Sociological Review* 48, no. 2 (1983): 147–160.

Epley, Nicholas, Adam Waytz, and John T. Cacioppo. "On Seeing Human: A Three-Factor Theory of Anthropomorphism." *Psychological Review* 114, no. 4 (2007): 864–886.

Goodhart, Charles A. E. "Problems of Monetary Management: The U.K. Experience." In *Papers in Monetary Economics*, vol. 1. Reserve Bank of Australia, 1975.

Meyer, John W., and Brian Rowan. "Institutionalized Organizations: Formal Structure as Myth and Ceremony." *American Journal of Sociology* 83, no. 2 (1977): 340–363.

Nass, Clifford, and Youngme Moon. "Machines and Mindlessness: Social Responses to Computers." *Journal of Social Issues* 56, no. 1 (2000): 81–103.

Reeves, Byron, and Clifford Nass. *The Media Equation: How People Treat Computers, Television, and New Media Like Real People and Places*. Cambridge: Cambridge University Press, 1996.

OpenAI. "GPT-4o System Card." August 8, 2024. <https://openai.com/index/gpt-4o-system-card/>

Waytz, Adam, John Cacioppo, and Nicholas Epley. "Who Sees Human? The Stability and Importance of Individual Differences in Anthropomorphism." *Perspectives on Psychological Science* 5, no. 3 (2010): 219–232.

Additional references may be added in a future journal-submission version addressing institutional epistemology, speech-act theory, philosophy of technology, AI governance communication, and human-computer interaction. The present version is finalized for archival deposit.