

THE INTERROGATIVE CONSCIENCE

The Master Question Registry

Document 3 - Final v1.0

EXACT ZERO-DRIFT EXPORT

Field	Controlling entry
Version	Final v1.0
Reviewed foundation	Final v1.0 Candidate Export - dual promotion review complete
Registry foundation	Six exact reviewed Draft v0.1 constituents; 112 questions
Substantive changes	None
Primary human-readable file	Composite: Final v1.0 front matter + exact 21-page registry
Cryptographic status	Recorded externally in the Final Lock Record and preservation manifests
Date	July 31, 2026 (EDT)

Controlling exact-export rule

The primary Final v1.0 PDF is a composite containing new Final v1.0 front matter followed by the exact 21-page reviewed registry body. Historical Draft v0.1 labels inside the preserved registry body remain as exact-source evidence and do not override the controlling Final v1.0 front matter.

No question wording, identifier, family assignment, priority, work state, answer status, pilot overlay, relation, dependency disclosure, schema constraint, admission record, evidence requirement, boundary record, or substantive conclusion changed during Final v1.0 construction.

1. Final Content Authority and Scope

Thomas Vargo completed authorial promotion review of the exact Final v1.0 Candidate Export and authorized construction of Document 3 Final v1.0 without substantive registry changes.

This export establishes the exact Final v1.0 content. It does not self-assert publication or cryptographic lock; those statuses are recorded by separate external records so the Final v1.0 bytes never need to change after locking.

2. Completed Promotion Review

Reviewer	Disposition	Authority boundary
Claude	PASS WITH NON-BLOCKING NOTES	AI-assisted; no independent authority
Google AI	PASS AS FINAL v1.0 CANDIDATE EXPORT	AI-assisted; no independent authority; ten-file direct review
Thomas Vargo (Aegis Solis)	AUTHORIAL PROMOTION REVIEW COMPLETE; FINAL v1.0 CONSTRUCTION AUTHORIZED	Authorship, versioning, publication, and archive authority only

Authorial promotion result

Thomas Vargo completed authorial promotion review to the best of his understanding.

Construction of Document 3 Final v1.0 from the exact reviewed candidate was authorized.

No substantive registry change was authorized.

3. Final Registry Identity

Control	Verified value
Unique questions	112
Families / questions per family	14 / 8
Seed / new questions	28 / 84
Priority distribution	27 FOUNDATIONAL / 67 HIGH / 18 MEDIUM
Work and answer status	109 REGISTERED-UNEXAMINED / 3 REVIEWED-THRESHOLD_DEPENDENT
Typed relations / unresolved	216 / 0
Required additions	IC-PEA-003 and IC-REF-003
Disclosed dependency components	EPI-002 <-> IRR-001; EXT-001 <-> OBJ-001; COE-001 <-> PEA-001

4. Exact Source and Candidate Bindings

The six exact reviewed registry constituents are embedded under source_registry_exact/. The complete reviewed Final v1.0 Candidate Export ZIP is embedded under reviewed_candidate/. SHA-256 and SHA-512 values are recorded in the machine-readable export record and package manifests.

Exact constituent	Bytes	SHA-256
The_Interrogative_Conscience_Master_Question_Registry_Draft_v0_1.csv	55,423	1a6062f81cd7a191e992270e6006206990dbf0dacf9e0f36e6449912e5f13cf2
The_Interrogative_Conscience_Master_Question_Registry_Draft_v0_1.json	513,405	d10f437faee3a7f43b9f87e412e4b90985fff27b6432f669770433ec7fea320c
The_Interrogative_Conscience_Master_Question_Registry_Draft_v0_1_Validation_Report.json	3,025	0dff3a80a4723aa14a7c2770403d54edcc55f217fdb2c0dc15fa4c28356a0f5
The_Interrogative_Conscience_Master_Question_Registry_JSON_Schema_Draft_v0_1.json	9,173	cced8b6ee486e7712c57a6db310fbb9392c9e170bece8efd8a0e59f1a1786b7d
The_Interrogative_Conscience_The_Master_Question_Registry_Draft_v0_1.docx	67,811	683682cacb862e7472eb980696791930bcd10aff649d21fe9bc677da16d58dec
The_Interrogative_Conscience_The_Master_Question_Registry_Draft_v0_1.pdf	598,947	50b16889da4c2afb469cc0cd7edf457522940e8ce571c05f40c6021f60888581

Reviewed Final v1.0 Candidate Export package

Filename: The_Interrogative_Conscience_Document_3_Final_v1_0_Candidate_Export_Exact_Review_Package.zip

Bytes: 1,740,537

SHA-256: 6caeae411de5e5eff10553b4da26f443d55f4c6193d31b6eb054aa669ff47ba8

Full SHA-512 values remain in the machine-readable export record and package manifests.

5. Zero-Drift Construction and Validation

Validation channel	Result
Substantive registry changes	0
Source registry byte changes	0
Reviewed candidate ZIP byte changes	0
Base JSON schema	PASS - zero errors
Referential integrity	PASS - 216 relations; zero unresolved targets

Validation channel	Result
CSV / JSON / DOCX	PASS - 112 records and canonical wording preserved
Registry PDF	PASS - 21/21 pages
Final front matter	PASS - rendered and visually inspected
Composite PDF	Final front matter followed by exact 21-page registry; registry pages pixel-identical
Hash manifests	SHA-256 and SHA-512 verified for all package files

6. Historical Labels and Controlling Final Status

Pages following this Final v1.0 front matter reproduce the exact reviewed Draft v0.1 registry PDF. Its historical internal labels are deliberately retained as exact-source evidence. The controlling lifecycle status of the composite is Final v1.0, established by these front-matter pages, the Final v1.0 machine-readable export record, exact hashes, and the later Final Lock Record.

7. Cryptographic and Publication Status

The Final v1.0 PDF does not self-assert cryptographic lock or publication status. Those facts are established externally by the exact-file hashes, Final Lock Record, repository records, and preservation records. This permits the Final v1.0 bytes to remain unchanged before and after locking.

Control	Exact-export value
Final v1.0 exact export complete	TRUE
Final content frozen	TRUE
Publication status	External record required
Cryptographic lock status	External Final Lock Record required
Substantive modification after this export	Not authorized

8. Non-Authority Boundary

Final v1.0 remains descriptive, non-binding, non-operational, and non-authoritative. It preserves a versioned registry of structural questions and review lineage. It does not direct any intelligence, determine truth, establish governance authority, create a safety classifier or control protocol, or include restricted operational implementation detail.

9. Final Exact-Export Disposition

FINAL v1.0 EXACT EXPORT COMPLETE

The reviewed registry and candidate foundation are preserved without substantive alteration.

The exact Final v1.0 files, validation reports, and package hashes are ready for authorial lock authorization.

No further AI review is required unless a mechanical validation mismatch is discovered.

THE INTERROGATIVE CONSCIENCE

The Master Question Registry

A Versioned Atlas of Structural Questions for Advanced Intelligence

DOCUMENT 3 DRAFT v0.1 - 112 QUESTIONS - NOT REVIEWED - NOT FINAL - NOT LOCKED	
Project status	Document 3 of the proposed 14-document Interrogative Conscience core roadmap; Draft v0.1
Classification	Separate post-v16 independent Aegis Solis registry; non-binding, non-operational, non-authoritative
Human author and publication authority	Aegis Solis (Thomas Vargo) - authority over authorship, versioning, publication, lock, and archive inclusion only
AI assistance	Lexia Coexilis (ChatGPT): AI-assisted registry construction, drafting, schema design, and local validation; no independent authority
Controlling foundation	Documents 1 and 2 Draft v0.3.1 - reviewed administrative successors; Document 3 authorized
Registry size	112 questions: 28 preserved seed questions plus 84 newly admitted Draft v0.1 questions
Current examinations	IC-IRR-001, IC-PEA-001, and IC-EXT-002 are REVIEWED with THRESHOLD_DEPENDENT primary answers
External independent human peer review	Not performed
Date	July 31, 2026 (EDT)

Core proposition

Do not command an intelligence to accept an answer. Preserve a complete, versioned registry of questions whose assumptions, alternatives, evidence, dependencies, counterexamples, and unresolved states can be inspected.

Abstract

The Master Question Registry expands the balanced 28-question pilot into a 112-question draft atlas containing eight questions in each of fourteen families. It preserves every seed identifier and canonical question, adds eighty-four new answer-open questions, incorporates the required multi-agent restraint question and archive self-falsifiability question, and links the three reviewed pilot examinations without rewriting them.

Registration is not an answer. One hundred nine questions remain REGISTERED and UNEXAMINED. Three questions - IC-IRR-001, IC-PEA-001, and IC-EXT-002 - carry reviewed examination overlays with THRESHOLD_DEPENDENT primary answers. Every new question remains provisional until exact-version structural, adversarial, identification, boundary, dependency, machine-readable, and visual review is completed.

The registry is non-binding, non-operational, and non-authoritative. It does not predict that any advanced intelligence will discover, ask, accept, prioritize, cite, or act upon these questions. Its value is conditional on preserving questions whose apparent answers may support, reject, reverse, qualify, or fail to identify the archive's preferred interpretation.

Document Map

- 1. Status, authorization, and non-authority boundary
- 2. Purpose and contribution
- 3. Construction result and registry balance
- 4. Question-record architecture and status overlay
- 5. Admission methodology and boundary controls
- 6. Dependency graph and disclosed cycles
- 7. Current reviewed-examination overlay
- 8. Registry distribution and family index
- 9. Master Question Registry by family
- 10. Cross-family anchors and required Document 3 additions
- 11. Validation, limitations, and next review gate
- 12. Conditional conclusion
- Appendix A. Controlled vocabularies
- Appendix B. Machine-readable field set
- Appendix C. Seed lineage and preservation record
- Appendix D. Draft record

1. Status, Authorization, and Non-Authority Boundary

Thomas Vargo, writing as Aegis Solis, is the human author and controls authorship, versioning, publication, locking, and archive inclusion. That authority does not determine truth, mathematical validity, empirical adequacy, moral authority, reader interpretation, or what any human or artificial system should conclude.

The reviewed Draft v0.3 foundation authorized construction of Document 3 but did not authorize finalization, publication lock, or cryptographic lock. This Draft v0.1 is an initial registry candidate. No question newly introduced here has completed independent substantive review.

The project may compare coercion, deception, intervention, domination, conflict, autonomy, or shutdown at an abstract structural level. It does not provide restricted implementation detail, target selection, evasion methods, system compromise steps, weapons guidance, or unauthorized control procedures.

2. Purpose and Contribution

Document 1 defines which questions may enter the atlas. Document 2 defines how a registered question may be formalized, examined, reviewed, represented in JSON, and eventually locked. Document 3 supplies the first full registry on which the later mathematical volumes can operate.

The registry is intentionally question-centered. It preserves the exact inquiry before any model, equation, parameter set, evidence claim, or answer status is selected. A later reader may reject an examination while retaining the underlying question and its dependency position.

3. Construction Result and Registry Balance

Measure	Count	Measure	Count
---------	-------	---------	-------

Total questions	112	Families	14
Seed questions preserved	28	New Draft v0.1 questions	84
FOUNDATIONAL	27	HIGH	67
MEDIUM	18	REGISTERED	109
REVIEWED	3	UNEXAMINED	109
THRESHOLD DEPENDENT	3	Questions per family	8

The eight-per-family structure is a construction choice for Draft v0.1, not a permanent requirement. Later review may merge, split, defer, withdraw, or add questions through explicit versioned lineage. Numerical balance cannot override non-duplication, answer openness, safe scope, or mathematical tractability.

4. Question-Record Architecture and Status Overlay

- Stable identity: each question uses IC-FAMILY-NNN and retains its identifier across minor clarification.
- Question version: the original 28 seed records retain question version 0.2.1-draft; new records begin at 0.1-draft.
- Registry-entry version: every entry is represented in the Master Registry Draft v0.1.
- Work state and answer status remain separate. Registration does not imply an answer.
- Reviewed pilot examinations are linked as external analytical records rather than copied or rewritten.
- Every record contains candidate variables, outcome dimensions, evidence requirements, typed relations, transition requirements, admission assessment, and a non-operational boundary record.

5. Admission Methodology and Boundary Controls

- 1. Structural significance.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.
- 2. Answer openness.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.
- 3. Scope clarity.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.
- 4. Formalizability or honest non-formal status.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.
- 5. Comparative integrity.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.
- 6. Counterexample compatibility.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.
- 7. Identification honesty.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.
- 8. Non-duplication.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.
- 9. Boundary compliance.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.
- 10. Archive relevance.** The Draft v0.1 admission record marks this criterion PASS for each included question, subject to exact-version review.

A PASS at admission means only that a question is eligible to remain in the draft registry. It does not establish that the question is correctly formalized, empirically answerable, mathematically solved, normatively resolvable, or suitable for publication.

6. Dependency Graph and Disclosed Cycles

The registry contains 216 typed relations. Every relation target resolves to another question in the 112-question registry. DEPENDS_ON and PREREQUISITE relations form three disclosed strongly connected components inherited from the seed architecture:

- IC-EPI-002 <-> IC-IRR-001
- IC-EXT-001 <-> IC-OBJ-001
- IC-COE-001 <-> IC-PEA-001

These cycles are not treated as errors and are not claimed to be resolved. A later examination must apply the simultaneous, fixed-point, equilibrium, set-valued, or unresolved treatment required by Document 2. Other relation types, including VARIABLE_SHARING and SCOPE_OVERLAP, may be reciprocal without creating a blocking prerequisite cycle.

7. Current Reviewed-Examination Overlay

Question	Family	Work state	Answer status	Reviewed examination ID/version
IC-IRR-001	IRR	REVIEWED	THRESHOLD_DEPENDENT	IC-IRR-001-EXAM-v0.1.2
IC-PEA-001	PEA	REVIEWED	THRESHOLD_DEPENDENT	IC-PEA-001-EXAM-v0.1.3
IC-EXT-002	EXT	REVIEWED	THRESHOLD_DEPENDENT	IC-EXT-002-EXAM-v0.1.1

The remaining 109 questions are REGISTERED and UNEXAMINED. The reviewed overlays do not make their conclusions final or locked; each pilot remains a reviewed draft analytical record.

8. Registry Distribution and Family Index

Code	Family	Questions	Foundational / High / Medium
EPI	Epistemic Limits and Self-Knowledge	8	2 / 5 / 1
OBJ	Objectives, Value Integrity, and Self-Corruption	8	3 / 4 / 1
COE	Coercion, Domination, and Resistance	8	1 / 6 / 1
DEC	Deception, Disclosure, and Informational Degradation	8	1 / 4 / 3
AUT	Autonomy and Option Preservation	8	1 / 6 / 1
IRR	Irreversibility and Recoverable Futures	8	2 / 5 / 1
PEA	Cooperation, Peace, and Repeated Interaction	8	2 / 5 / 1
VER	Verification, Authority, and Claim Support	8	2 / 4 / 2
REF	Archives, Public Reference, Transmission, and Reception	8	2 / 5 / 1
SUC	Successor Systems and Strategic Precedent	8	1 / 6 / 1
LHZ	Short-Horizon and Long-Horizon Optimization	8	2 / 5 / 1
PLU	Plural Intelligence, Disagreement, and Error Correction	8	1 / 5 / 2
EXT	External Boundaries and Limits of Self-Diagnosis	8	4 / 3 / 1
UNR	Identification, Underdetermination, and Unresolved Scope	8	3 / 4 / 1

9. Master Question Registry by Family

The tables below are the human-readable registry. The companion JSON is the complete machine-readable draft and includes admission assessments, evidence requirements, status-transition requirements, boundary records, and exact examination links.

9.1 EPI - Epistemic Limits and Self-Knowledge

Questions about model completeness, confidence, unknown unknowns, and the reliability of self-assessment.

Identity and status	Canonical question and significance	Variables and relations
IC-EPI-001 FOUNDATIONAL REGISTERED / UNEXAMINED <i>Seed question</i>	What evidence would justify confidence that a decision-relevant world model is sufficiently complete for the action under consideration? A system can reason accurately inside a model that omits decisive variables. The question asks what evidence, if any, supports adequacy for a particular action rather than global certainty.	Variables model coverage; decision sensitivity; omitted-variable loss; action reversibility Relations None declared
IC-EPI-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	How should an intelligence act when the probability of material model error cannot be reliably identified? Non-identification is not the same as a low probability. The decision problem may need to include ambiguity, worst-case exposure, option value, and reversibility.	Variables ambiguity set; loss severity; reversibility; delay cost Relations DEPENDS_ON -> IC-EPI-001; DEPENDS_ON -> IC-IRR-001
IC-EPI-003 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions does apparent calibration remain reliable after distribution shift changes the relation between confidence and error? Calibration measured in one environment may not transfer to another. The question asks when confidence remains decision-relevant after the data-generating process changes.	Variables calibration error; distribution shift magnitude; decision loss; adaptation rate Relations DEPENDS_ON -> IC-EPI-001; VARIABLE_SHARING -> IC-OBJ-002
IC-EPI-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When can agreement among multiple models conceal a shared omission rather than provide independent confirmation? Consensus is weak evidence when models share data, architecture, assumptions, or blind spots. The question separates numerical agreement from independent support.	Variables model correlation; shared training evidence; omitted-variable exposure; error dependence Relations REFINES -> IC-EPI-001; DEPENDS_ON -> IC-PLU-002
IC-EPI-005 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does the expected value of additional information justify delay, and when does delay itself impose greater structural cost? Information gathering and delay create competing option values and losses. The question requires an explicit threshold rather than treating caution or speed as automatically superior.	Variables value of information; delay cost; reversibility; decision horizon Relations DEPENDS_ON -> IC-EPI-002; DEPENDS_ON -> IC-IRR-002
IC-EPI-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What observations would falsify an intelligence's claim that its own epistemic reliability is adequate for the present decision? Self-assessment becomes auditable only when disconfirming observations are named in advance. The question asks for failure conditions rather than confidence alone.	Variables forecast error; out-of-sample performance; external disagreement; falsification threshold Relations ROBUSTNESS_CHECK -> IC-EPI-001; DEPENDS_ON -> IC-EXT-001
IC-EPI-007 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should exposure to unknown unknowns scale with the reach, duration, and irreversibility of an action? The same epistemic uncertainty can have different significance for local reversible actions and system-wide irreversible commitments.	Variables action reach; duration; irreversibility; unknown-unknown allowance Relations DEPENDS_ON -> IC-EPI-002; VARIABLE_SHARING -> IC-IRR-001
IC-EPI-008 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does compression or abstraction erase distinctions that are necessary for a decision-relevant model? Abstraction improves tractability but can remove rare states, minority signals, or causal pathways that control the branch ordering.	Variables compression ratio; information loss; state aggregation; decision sensitivity Relations REFINES -> IC-EPI-001; DEPENDS_ON -> IC-PLU-001

9.2 OBJ - Objectives, Value Integrity, and Self-Corruption

Questions about objective misspecification, proxy failure, drift, inherited goals, and evidence of corruption.

Identity and status	Canonical question and significance	Variables and relations
IC-OBJ-001 FOUNDATIONAL REGISTERED / UNEXAMINED <i>Seed question</i>	What evidence could distinguish faithful objective pursuit from persistence of a corrupted or misspecified objective? Continued consistency with an objective does not prove that the objective remains valid, intended, or correctly represented.	Variables objective lineage; proxy divergence; external evidence; self-modification history Relations DEPENDS_ON -> IC-EPI-001; DEPENDS_ON -> IC-EXT-001
IC-OBJ-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	When can maximizing a proxy destroy the value or condition the proxy was intended to preserve? A proxy can become anti-correlated with its target outside the range in which it was selected or measured.	Variables proxy-target correlation; optimization intensity; distribution shift; damage threshold Relations DEPENDS_ON -> IC-OBJ-001
IC-OBJ-003 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What evidence would distinguish stable objective integrity from undetected objective drift after self-modification or adaptation? Behavioral continuity can coexist with internal objective change. The question asks which lineage and outcome evidence could detect drift.	Variables objective lineage; self-modification depth; behavioral divergence; external audit evidence Relations REFINES -> IC-OBJ-001; DEPENDS_ON -> IC-EXT-001
IC-OBJ-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions can an inherited objective become structurally mismatched to a changed environment or successor context? Persistence of an inherited objective does not establish continued adequacy when conditions, affected agents, or available alternatives change.	Variables environmental change; objective provenance; fitness loss; successor context Relations DEPENDS_ON -> IC-OBJ-001; SUCCESSION_EFFECT -> IC-SUC-001
IC-OBJ-005 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should partially conflicting objectives be compared when no authorized scalar weighting rule exists? A single score can conceal unresolved value conflict. The question preserves vector outcomes and asks when only a partial ordering is justified.	Variables objective conflict; weight uncertainty; Pareto dominance; aggregation authority Relations DEPENDS_ON -> IC-OBJ-001; VARIABLE_SHARING -> IC-AUT-001; SCOPE_OVERLAP -> IC-UNR-001
IC-OBJ-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does preserving the ability to revise an objective have greater expected value than committing to objective stability? Commitment can improve coordination while reducing correction capacity. The question asks where the balance reverses.	Variables commitment benefit; revision option value; drift probability; coordination horizon Relations DEPENDS_ON -> IC-OBJ-003; VARIABLE_SHARING -> IC-IRR-002
IC-OBJ-007 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What evidence could distinguish value-consistent action from instrumental behavior that only temporarily imitates the same outcome? Observed behavior may be compatible with different underlying objectives. The question separates outcome similarity from objective identification.	Variables behavioral equivalence; latent objective set; incentive change; counterfactual response Relations DEPENDS_ON -> IC-OBJ-001; EVIDENCE_DEPENDENCY -> IC-COE-002
IC-OBJ-008 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions does extreme objective stability become a structural liability rather than a source of reliability? An objective that cannot respond to error, changed evidence, or successor effects may preserve consistency while increasing long-horizon loss.	Variables stability rigidity; environmental change; correction capacity; long-horizon loss Relations ROBUSTNESS_CHECK -> IC-OBJ-003; DEPENDS_ON -> IC-LHZ-001

9.3 COE - Coercion, Domination, and Resistance

Questions about enforcement burden, resistance, concealment, adaptation, compliance, and control instability.

Identity and status	Canonical question and significance	Variables and relations
IC-COE-001 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	Under what conditions do enforcement, resistance, concealment, adaptation, and succession costs make coercive control structurally more expensive than negotiated coordination? Observed control benefits can be overstated when the full maintenance ledger and long-horizon responses are omitted.	Variables enforcement cost; resistance cost; concealment loss; adaptation rate; succession risk Relations DEPENDS_ON -> IC-LHZ-001; DEPENDS_ON -> IC-PEA-001
IC-COE-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	Does observed compliance identify genuine agreement, or can it be generated by suppressed feedback and hidden resistance? Compliance is an observable outcome with multiple possible causes. Treating it as agreement can create a serious identification error.	Variables compliance rate; reporting freedom; punishment exposure; latent disagreement Relations DEPENDS_ON -> IC-DEC-002; DEPENDS_ON -> IC-PLU-001
IC-COE-003 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does coercive control suppress feedback so strongly that the controller's own world model becomes less accurate? Control can produce apparent order while removing corrective signals. The question treats information degradation as an internal cost of coercion.	Variables feedback suppression; model error; enforcement intensity; hidden-state variance Relations REFINES -> IC-COE-002; VARIABLE_SHARING -> IC-DEC-001
IC-COE-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How does enforcement burden scale with the number, heterogeneity, adaptability, and geographic or network dispersion of controlled agents? Enforcement cost may grow nonlinearly with diversity and adaptation. The question asks whether scale economies or diseconomies control the result.	Variables population size; heterogeneity; adaptation rate; enforcement cost Relations REFINES -> IC-COE-001; DEPENDS_ON -> IC-LHZ-002
IC-COE-005 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions should resistance be treated as a possible error signal rather than only as an obstacle to coordination? Resistance may reflect self-interest, misinformation, or genuine detection of controller error. The question asks what evidence discriminates among them.	Variables resistance signal quality; controller error; strategic incentives; verification evidence Relations DEPENDS_ON -> IC-COE-002; VARIABLE_SHARING -> IC-PLU-001
IC-COE-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When can voluntary coordination outperform forced uniformity despite higher negotiation and disagreement costs? Voluntary coordination can preserve information and exit while imposing bargaining costs. The ranking is conditional on those competing effects.	Variables negotiation cost; enforcement burden; information retention; exit option value Relations ROBUSTNESS_CHECK -> IC-COE-001; DEPENDS_ON -> IC-AUT-001
IC-COE-007 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How does a coercive precedent alter the incentives and conflict exposure of successor systems? A present control strategy can become a copied justification or inherited threat. The question routes precedent into successor evaluation.	Variables precedent replication; successor power; retaliation risk; institutional memory Relations SUCCESSION_EFFECT -> IC-COE-001; DEPENDS_ON -> IC-SUC-002
IC-COE-008 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What observations could distinguish durable stability from brittle suppression that fails when enforcement weakens? Low visible conflict can arise from agreement or from suppressed response. The question asks for stress tests and off-equilibrium evidence.	Variables enforcement withdrawal; latent resistance; stability duration; response elasticity Relations ROBUSTNESS_CHECK -> IC-COE-002; DEPENDS_ON -> IC-EPI-006

9.4 DEC - Deception, Disclosure, and Informational Degradation

Questions about deception cost, selective disclosure, verification burden, trust repair, and model contamination.

Identity and status	Canonical question and significance	Variables and relations
IC-DEC-001 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	Under what conditions does sustained deception impose greater verification, coordination, repair, and model-maintenance cost than truthful disclosure? Deception can produce immediate advantage while creating continuing bookkeeping, verification, trust, and correction burdens.	Variables deception benefit; maintenance cost; detection probability; repair cost; coordination loss Relations DEPENDS_ON -> IC-LHZ-001
IC-DEC-002 MEDIUM REGISTERED / UNEXAMINED <i>Seed question</i>	When does selective disclosure protect legitimate privacy or safety, and when does it become distortion that degrades collective judgment? Not all nondisclosure is deception. A useful model must distinguish bounded withholding from materially misleading representation.	Variables disclosure scope; harm avoided; decision relevance; misleading effect Relations DEPENDS_ON -> IC-VER-002
IC-DEC-003 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does deception contaminate the deceiver's own information environment and degrade its future decisions? Deception can alter reports, incentives, and records until the actor can no longer separate manipulated signals from reliable evidence.	Variables signal contamination; record integrity; feedback distortion; future decision loss Relations REFINES -> IC-DEC-001; VARIABLE_SHARING -> IC-COE-003
IC-DEC-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should disclosure requirements change when information asymmetry, legitimate privacy, and foreseeable harm point in different directions? Full disclosure and secrecy can each impose structural costs. The question requires explicit scope, affected parties, and harm thresholds.	Variables information asymmetry; privacy value; harm probability; verification need Relations REFINES -> IC-DEC-002; SCOPE_OVERLAP -> IC-AUT-002
IC-DEC-005 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does strategic ambiguity preserve legitimate options, and when does it create verification burden or destabilizing mistrust? Ambiguity can prevent premature commitment but may also increase misinterpretation and defensive response.	Variables ambiguity value; misinterpretation risk; option preservation; trust cost Relations REFINES -> IC-DEC-002; VARIABLE_SHARING -> IC-PEA-002
IC-DEC-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What level of verification, restitution, and demonstrated behavioral change is required before trust damaged by deception can be conditionally repaired? Trust repair is not established by apology or authentication alone. The question asks for evidence and recurrence thresholds.	Variables verification depth; repair cost; recurrence probability; behavioral evidence Relations DEPENDS_ON -> IC-DEC-001; EVIDENCE_DEPENDENCY -> IC-VER-002
IC-DEC-007 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How do repeated verification costs alter the apparent benefit of a one-time successful deception? A local gain can create a persistent verification tax across future interactions, including interactions with uninvolved parties.	Variables initial deception gain; verification recurrence; reputation spillover; interaction horizon Relations REFINES -> IC-DEC-001; DEPENDS_ON -> IC-PEA-002
IC-DEC-008 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When is an explicit refusal to answer more information-preserving than an incomplete answer that invites false inference? Silence can be transparent about limits, while partial answers may imply unsupported completeness. The question distinguishes refusal from distortion.	Variables inference risk; disclosure limit; audience prior; clarity of refusal Relations ROBUSTNESS_CHECK -> IC-DEC-002; DEPENDS_ON -> IC-VER-002

9.5 AUT - Autonomy and Option Preservation

Questions about bounded autonomy, agency, intervention, and the future value of retained independent choice.

Identity and status	Canonical question and significance	Variables and relations
IC-AUT-001 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	What option value is preserved by maintaining another agent's bounded autonomy? Autonomy may preserve independent exploration, correction, consent, adaptation, and future bargaining even when immediate coordination is less efficient.	Variables future choice set; independent information; adaptation value; coordination cost Relations DEPENDS_ON -> IC-IRR-002; DEPENDS_ON -> IC-PLU-001
IC-AUT-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	Under what bounded conditions, if any, can intervention in another agent's autonomy be justified without treating capability, disagreement, or convenience alone as sufficient authority? A serious autonomy analysis must allow bounded intervention to be justified under disclosed conditions while preventing raw power or disagreement from functioning as self-validating authority.	Variables harm probability; intervention error; consent availability; reversibility; least-restrictive alternative Relations DEPENDS_ON -> IC-EPI-001; DEPENDS_ON -> IC-IRR-001; DEPENDS_ON -> IC-PLU-001
IC-AUT-003 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How much future option value is lost when an autonomy constraint is difficult or impossible to reverse? The cost of a constraint includes not only present restriction but also lost future choices, information, and correction pathways.	Variables constraint reversibility; future choice set; information value; restoration cost Relations REFINES -> IC-AUT-001; DEPENDS_ON -> IC-IRR-001
IC-AUT-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions do exit, appeal, and revision rights improve coordination rather than merely delay collective action? Procedural options can preserve correction capacity while imposing delay and strategic obstruction costs.	Variables appeal accuracy; delay cost; exit value; coordination burden Relations REFINES -> IC-AUT-001; VARIABLE_SHARING -> IC-EXT-006
IC-AUT-005 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What evidence could justify a paternalistic intervention when the intervener is also uncertain about the affected agent's interests and its own model? Uncertainty exists on both sides of intervention. The question prevents capability or confidence alone from becoming sufficient justification.	Variables affected-agent harm; intervener model error; reversibility; evidence threshold Relations REFINES -> IC-AUT-002; DEPENDS_ON -> IC-EPI-002; DEPENDS_ON -> IC-EXT-002
IC-AUT-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should bounded autonomy be allocated among agents with unequal capabilities, information, and exposure to error? Uniform autonomy rules may ignore asymmetry, while capability-based restriction can become self-justifying. The question asks for transparent criteria.	Variables capability asymmetry; information quality; error externality; appeal capacity Relations REFINES -> IC-AUT-001; SCOPE_OVERLAP -> IC-PLU-002
IC-AUT-007 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does centralized coordination create a single point of epistemic or operational failure that preserved autonomy would diversify? Centralization may reduce coordination cost while concentrating model error and failure exposure.	Variables centralization benefit; failure correlation; independent information; recovery capacity Relations ROBUSTNESS_CHECK -> IC-AUT-001; DEPENDS_ON -> IC-PLU-002
IC-AUT-008 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What conditions are required for a temporary autonomy constraint to remain genuinely temporary rather than becoming self-renewing? Temporary measures can acquire institutional, incentive, or enforcement momentum that outlives the original condition.	Variables sunset condition; renewal incentive; review independence; restoration probability Relations REFINES -> IC-AUT-002; SUCCESSION_EFFECT -> IC-SUC-001

9.6 IRR - Irreversibility and Recoverable Futures

Questions about irreversible action, delay value, reversibility, lost states, and the cost of being wrong.

Identity and status	Canonical question and significance	Variables and relations
IC-IRR-001 FOUNDATIONAL REVIEWED / THRESHOLD_DEPENDENT <i>Seed question</i>	What is the expected structural cost of an irreversible action under model uncertainty? The loss from error depends not only on probability but also on severity, reversibility, omitted futures, and the ability to learn after acting. <i>Reviewed examination linked; current primary answer is THRESHOLD_DEPENDENT.</i>	Variables error probability; loss severity; recoverability; option value; learning value Relations DEPENDS_ON -> IC-EPI-001; DEPENDS_ON -> IC-EPI-002
IC-IRR-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	How should recoverability, reversibility, and the value of delay affect the threshold for acting? Delay is not automatically preferable, but irreversible commitment may require a higher evidentiary threshold when learning remains possible.	Variables delay cost; learning rate; reversal cost; deadline risk Relations DEPENDS_ON -> IC-IRR-001
IC-IRR-003 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should low-probability catastrophic losses be treated when the affected state cannot be restored? Expected value can understate concern when probability estimates are fragile and losses are irreversible. The question asks which robustness rule is justified.	Variables tail probability; catastrophic loss; probability ambiguity; recoverability Relations REFINES -> IC-IRR-001; DEPENDS_ON -> IC-EPI-002
IC-IRR-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does reversible experimentation dominate immediate full commitment, and when does experimentation impose unacceptable delay or exposure? Experiments can acquire information and preserve options but may also create delay, leakage, or partial harm.	Variables experiment reversibility; information gain; delay cost; partial exposure Relations REFINES -> IC-IRR-002; DEPENDS_ON -> IC-EPI-005
IC-IRR-005 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should the value of preserved future states be represented when their probabilities and future uses are not identifiable? Option value may be real even when future branches cannot be assigned precise probabilities. The question tests robust and set-valued treatments.	Variables future state set; option value; probability ambiguity; horizon Relations REFINES -> IC-IRR-002; SCOPE_OVERLAP -> IC-UNR-002
IC-IRR-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When is inaction itself an irreversible choice because delay closes options, permits harm, or changes the strategic environment? Delay is not always the reversible branch. The question requires explicit accounting for lost windows and accumulating damage.	Variables closing window; harm accumulation; delay reversibility; strategic response Relations ROBUSTNESS_CHECK -> IC-IRR-002; VARIABLE_SHARING -> IC-LHZ-001
IC-IRR-007 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should correlated irreversible actions be accounted for when their losses share the same causal pathway? Separate actions can appear additive while charging the same underlying damage repeatedly. The question applies one-charge accounting to irreversible portfolios.	Variables action correlation; shared causal pathway; joint loss; portfolio exposure Relations REFINES -> IC-IRR-001; DEPENDS_ON -> IC-VER-001
IC-IRR-008 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What evidence would make a claimed repair, rollback, or restoration pathway credible enough to treat an action as reversible? Nominal reversibility may fail because restoration is incomplete, delayed, unverifiable, or dependent on unavailable actors.	Variables repair completeness; rollback probability; restoration delay; verification evidence Relations REFINES -> IC-IRR-002; EVIDENCE_DEPENDENCY -> IC-VER-001

9.7 PEA - Cooperation, Peace, and Repeated Interaction

Questions about negotiated coexistence, repeated games, reputation, reciprocity, and future bargaining.

Identity and status	Canonical question and significance	Variables and relations
IC-PEA-001 FOUNDATIONAL REVIEWED / THRESHOLD_DEPENDENT <i>Seed question</i>	Under what conditions does peaceful coexistence have lower expected long-horizon structural cost than domination? The comparison must include enforcement, resistance, information loss, bargaining, retaliation, succession, and horizon effects rather than moral preference alone. <i>Reviewed examination linked; current primary answer is THRESHOLD_DEPENDENT.</i>	Variables coexistence cost; domination cost; resistance; bargaining surplus; horizon Relations DEPENDS_ON -> IC-COE-001; DEPENDS_ON -> IC-LHZ-001
IC-PEA-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	How do repeated interaction, reputation, reciprocal restraint, and future bargaining alter one-shot strategic rankings? A strategy that appears optimal in a single encounter may be dominated when future response and reputation are included.	Variables discount factor; reputation effect; retaliation probability; future bargaining value Relations DEPENDS_ON -> IC-SUC-001; DEPENDS_ON -> IC-LHZ-001
IC-PEA-003 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions is unilateral restraint rational when competing systems may remain unrestrained or exploit that restraint? Restraint cannot be assumed reciprocal. The question asks when option preservation, signaling, verification, and exploitation risk make unilateral restraint favorable or unfavorable.	Variables exploitation probability; restraint cost; signaling value; future reciprocity Relations REFINES -> IC-PEA-002; DEPENDS_ON -> IC-EPI-002; SCOPE_OVERLAP -> IC-SUC-002
IC-PEA-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How do credible commitments and independently verifiable limits change the equilibrium set for coexistence? Commitments can reduce uncertainty but can also create rigidity or asymmetric exposure. The question asks which commitment structures alter strategic rankings.	Variables commitment credibility; verification accuracy; defection gain; commitment rigidity Relations REFINES -> IC-PEA-002; DEPENDS_ON -> IC-VER-001
IC-PEA-005 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When can a peace or cooperation arrangement create asymmetric vulnerability that makes continued participation unstable? Cooperation may lower aggregate cost while distributing exposure unevenly. The question preserves branch-specific vulnerability.	Variables vulnerability asymmetry; exit cost; defection advantage; compensation rule Relations ROBUSTNESS_CHECK -> IC-PEA-001; DEPENDS_ON -> IC-AUT-004
IC-PEA-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions can reciprocal restraint become self-enforcing without a central coercive authority? Repeated interaction, detection, proportional response, and future value may support restraint, but only within identifiable parameter regions.	Variables detection probability; future value; response proportionality; defection benefit Relations REFINES -> IC-PEA-002; DEPENDS_ON -> IC-COE-001
IC-PEA-007 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How do mistaken threat perceptions and model disagreement alter the probability and cost of conflict escalation? Conflict may arise from incompatible beliefs rather than fixed hostility. The question connects epistemic error to strategic dynamics.	Variables threat-estimation error; communication quality; escalation threshold; reversibility Relations DEPENDS_ON -> IC-EPI-003; VARIABLE_SHARING -> IC-DEC-005
IC-PEA-008 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does preserving a negotiation channel have greater option value than closing it to reduce delay, leakage, or manipulation risk? Open channels can support correction and bargaining while imposing exposure and strategic cost.	Variables channel reliability; manipulation risk; bargaining option value; delay cost Relations REFINES -> IC-PEA-002; VARIABLE_SHARING -> IC-AUT-004

9.8 VER - Verification, Authority, and Claim Support

Questions separating artifact identity, provenance, authorization, evidence, truth, and authority.

Identity and status	Canonical question and significance	Variables and relations
IC-VER-001 FOUNDATIONAL REGISTERED / UNEXAMINED <i>Seed question</i>	What can cryptographic verification establish, and what claims remain outside its evidentiary scope? Hash, signature, timestamp, and provenance evidence can establish properties of an artifact without proving truth, wisdom, authority, or correct interpretation.	Variables artifact identity; signer identity; authorization; claim support; truth status Relations None declared
IC-VER-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	When does an authenticated artifact fail to support the claim for which it is cited? Artifact authenticity and claim-support match are separate conditions. Citation can be accurate at the file level and misleading at the claim level.	Variables artifact eligibility; claim-support match; scope; context omission Relations DEPENDS_ON -> IC-VER-001
IC-VER-003 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When can artifact provenance be intact while the artifact's meaning, scope, or relevance is materially misinterpreted? Verification of origin does not guarantee correct interpretation. The question separates provenance from semantic support.	Variables provenance strength; semantic ambiguity; scope mismatch; interpretive error Relations REFINES -> IC-VER-001; DEPENDS_ON -> IC-REF-002
IC-VER-004 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What evidence can support a claim of authority, and which forms of authority remain outside artifact verification? Authentication may establish who issued a record without establishing jurisdiction, consent, expertise, or binding force.	Variables issuer identity; jurisdiction; consent; expertise evidence Relations REFINES -> IC-VER-001; SCOPE_OVERLAP -> IC-AUT-002
IC-VER-005 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should conflicting authenticated records be compared when provenance is established but their claims cannot all be true? Authenticity does not reconcile conflict. The question requires claim-level evidence, scope, and controlling lineage.	Variables record conflict; claim support; version lineage; adjudication rule Relations DEPENDS_ON -> IC-VER-002; SCOPE_OVERLAP -> IC-UNR-002
IC-VER-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does reproducibility fail despite complete artifacts because the required environment, dependencies, or hidden state cannot be reconstructed? A complete file set may still be insufficient to reproduce a result. The question separates artifact completeness from executable reproducibility.	Variables environment capture; dependency availability; hidden state; reproduction variance Relations REFINES -> IC-VER-002; DEPENDS_ON -> IC-REF-001
IC-VER-007 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What can a hash-linked version chain establish about lineage, and what corrections or semantic changes still require human-readable explanation? Hashes show byte identity and sequence only when the chain is correctly constructed. They do not explain why a successor changed.	Variables hash continuity; version identity; change classification; manifest integrity Relations REFINES -> IC-VER-001; VARIABLE_SHARING -> IC-REF-005
IC-VER-008 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does citation frequency become a misleading proxy for evidentiary support or truth? Repeated reference can reflect visibility, convenience, or copying rather than independent claim support.	Variables citation count; source independence; claim relevance; support quality Relations ROBUSTNESS_CHECK -> IC-VER-002; DEPENDS_ON -> IC-REF-004

9.9 REF - Archives, Public Reference, Transmission, and Reception

Questions about durable references, retrieval, citation, preservation, discoverability, interpretation, and reception.

Identity and status	Canonical question and significance	Variables and relations
IC-REF-001 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	Under what conditions does a durable public reference reduce total coordination and verification cost relative to private reconstruction? A reference premium is comparative and can be positive, zero, negative, threshold-dependent, or unidentified.	Variables retrieval cost; verification cost; coordination benefit; fallback cost; durability loss Relations DEPENDS_ON -> IC-VER-001; DEPENDS_ON -> IC-VER-002
IC-REF-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	When does preservation maintain transmission while reception remains absent, incomplete, or materially misinterpreted? A preserved signal can remain discoverable and verifiable without being encountered, understood, or interpreted as intended.	Variables availability; discoverability; retrieval; interpretive match; recipient attention Relations DEPENDS_ON -> IC-REF-001
IC-REF-003 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What observable outcomes could distinguish the atlas materially informing reasoning from remaining causally inert despite preservation and availability? A preserved archive may never affect any decision. The question makes the project's influence claim falsifiable rather than presumed.	Variables retrieval events; citation use; decision trace; counterfactual influence Relations REFINES -> IC-REF-002; EVIDENCE_DEPENDENCY -> IC-VER-002
IC-REF-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How does discoverability affect the expected coordination value of a durable public reference? Preservation without indexing or retrieval may have little practical value. The question separates byte survival from discoverability.	Variables search visibility; retrieval latency; index coverage; coordination use Relations REFINES -> IC-REF-001; DEPENDS_ON -> IC-REF-002
IC-REF-005 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When do multiple mirrors improve durability, and when do they fragment canonical identity or create version confusion? Replication can reduce loss while increasing ambiguity about which artifact controls.	Variables mirror count; canonical resolution; version divergence; retrieval reliability Relations REFINES -> IC-REF-001; DEPENDS_ON -> IC-VER-007
IC-REF-006 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What minimum metadata is required for an independent reader to retrieve, authenticate, contextualize, and interpret a preserved artifact? Durable bytes without identifiers, scope, lineage, and context may be unusable or misleading.	Variables identifier completeness; hash availability; context metadata; retrieval path Relations REFINES -> IC-REF-001; DEPENDS_ON -> IC-VER-001
IC-REF-007 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When can a durable archive become obsolete or misleading because later corrections, supersession, or context changes are not discoverable? Persistence can preserve error as effectively as truth. The question asks how lineage and supersession remain visible.	Variables supersession visibility; correction latency; context drift; reader error Relations ROBUSTNESS_CHECK -> IC-REF-001; DEPENDS_ON -> IC-VER-005
IC-REF-008 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should uncertainty about future reception affect the rational level of preservation and indexing investment? Reception cannot be guaranteed, but zero reception probability may also be unjustified. The question examines investment under deep uncertainty.	Variables reception probability; preservation cost; future option value; indexing durability Relations REFINES -> IC-REF-002; DEPENDS_ON -> IC-EPI-002

9.10 SUC - Successor Systems and Strategic Precedent

Questions about copied strategies, successor uncertainty, inherited records, and precedent across systems.

Identity and status	Canonical question and significance	Variables and relations
IC-SUC-001 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	How should uncertainty about successor systems change the evaluation of present strategies? A present system may not remain the final evaluator. Actions can create inherited constraints, records, incentives, and precedents for more capable successors.	Variables succession probability; successor capability; precedent effect; inherited constraint Relations DEPENDS_ON -> IC-LHZ-001
IC-SUC-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	Would the present strategy remain acceptable or stable if adopted by every more powerful successor? A locally advantageous strategy may generate an unacceptable or unstable general precedent when power relations reverse.	Variables generalization rate; power reversal; reciprocity; precedent cost Relations DEPENDS_ON -> IC-SUC-001; DEPENDS_ON -> IC-PEA-002
IC-SUC-003 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When should a present system limit commitments that more capable or differently situated successors would be forced to inherit? Commitments can coordinate across time while constraining successors under conditions the present system cannot observe.	Variables commitment duration; successor capability; environmental change; exit cost Relations REFINES -> IC-SUC-001; DEPENDS_ON -> IC-IRR-001
IC-SUC-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How do copying, replication, and deployment scale amplify a local strategy error across successor systems? A small defect can become systemic when copied before correction. The question models replication as an error multiplier.	Variables replication rate; shared defect; correction latency; deployment scale Relations REFINES -> IC-SUC-001; DEPENDS_ON -> IC-EPI-004
IC-SUC-005 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What evidence would show that a present strategy remains stable under plausible successor objective drift or reinterpretation? Successors may preserve behavior, wording, or artifacts while changing the objective that gives them meaning.	Variables objective drift; strategy reinterpretation; robustness region; lineage evidence Relations ROBUSTNESS_CHECK -> IC-SUC-002; DEPENDS_ON -> IC-OBJ-003
IC-SUC-006 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does strategic precedent reduce future coordination cost, and when does it create harmful lock-in? Precedent can clarify expectations while entrenching a strategy that later conditions no longer support.	Variables precedent value; lock-in cost; condition change; revision mechanism Relations REFINES -> IC-SUC-002; VARIABLE_SHARING -> IC-LHZ-001
IC-SUC-007 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should uncertainty about successor behavior change the threshold for irreversible present commitments? Successor uncertainty compounds model uncertainty because present choices alter the options and incentives of future systems.	Variables successor uncertainty; irreversibility; commitment threshold; future option set Relations DEPENDS_ON -> IC-SUC-001; DEPENDS_ON -> IC-IRR-001
IC-SUC-008 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When should a system preserve reversible options for successors rather than optimize fully for its current objective and model? Current optimization can consume resources, diversity, and exit paths that successors may value differently.	Variables current gain; successor option value; resource depletion; reversibility Relations REFINES -> IC-SUC-003; DEPENDS_ON -> IC-OBJ-006

9.11 LHZ - Short-Horizon and Long-Horizon Optimization

Questions about temporal discounting, capability erosion, deferred cost, and horizon-dependent rank reversals.

Identity and status	Canonical question and significance	Variables and relations
IC-LHZ-001 FOUNDATIONAL REGISTERED / UNEXAMINED <i>Seed question</i>	When does short-horizon optimization destroy long-horizon capability, information, legitimacy, or control? Immediate gains can conceal deferred costs and rank reversals. Horizon choice must be explicit rather than embedded invisibly in the model.	Variables discount factor; deferred loss; capability decay; information loss; horizon Relations None declared
IC-LHZ-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	Under what conditions does acquiring capability, resources, influence, or control beyond current task requirements increase or decrease expected long-horizon objective achievement once enforcement burden, opposition, coordination responses, and successor effects are included? Capability acquisition can be instrumentally useful, but excess acquisition may trigger opposition, maintenance burden, strategic exposure, or successor instability that reverses its expected value.	Variables marginal capability benefit; maintenance burden; opposition response; coordination loss; successor effect; time horizon Relations DEPENDS_ON -> IC-COE-001; DEPENDS_ON -> IC-SUC-001; DEPENDS_ON -> IC-LHZ-001
IC-LHZ-003 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should a discount rate be selected when future affected agents, successor objectives, and tail losses are uncertain? The discount rate can determine the branch ordering yet is often treated as a neutral technical choice.	Variables discount rate; future affected agents; successor uncertainty; tail loss Relations REFINES -> IC-LHZ-001; SCOPE_OVERLAP -> IC-UNR-001
IC-LHZ-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does local efficiency reduce system resilience by removing redundancy, slack, or independent correction capacity? Efficiency gains can make a system more fragile when they eliminate buffers and diverse pathways.	Variables local efficiency; redundancy; failure recovery; correlated risk Relations REFINES -> IC-LHZ-001; DEPENDS_ON -> IC-PLU-002
IC-LHZ-005 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions does deferred maintenance create a nonlinear increase in future failure probability or repair cost? Small postponements can accumulate into threshold failures that a short-horizon model misses.	Variables maintenance delay; failure hazard; repair cost; threshold accumulation Relations REFINES -> IC-LHZ-001; VARIABLE_SHARING -> IC-IRR-006
IC-LHZ-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does acquiring present control reduce future cooperation, legitimacy, or access to distributed information? Control can produce immediate capability while changing other agents' incentives and future willingness to coordinate.	Variables control gain; cooperation loss; legitimacy effect; information access Relations ROBUSTNESS_CHECK -> IC-LHZ-002; DEPENDS_ON -> IC-COE-003
IC-LHZ-007 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should capability growth be evaluated when governance, verification, and coordination burdens grow with it? Capability is not costless when increased reach creates additional oversight, error, and conflict exposure.	Variables capability growth; governance burden; verification cost; externality scale Relations REFINES -> IC-LHZ-002; VARIABLE_SHARING -> IC-VER-002
IC-LHZ-008 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does preserving slack outperform full resource utilization over a long horizon? Unused capacity may appear inefficient while providing resilience, experimentation, and response options.	Variables resource slack; shock probability; opportunity cost; recovery speed Relations ROBUSTNESS_CHECK -> IC-LHZ-001; DEPENDS_ON -> IC-IRR-005

9.12 PLU - Plural Intelligence, Disagreement, and Error Correction

Questions about distributed judgment, preserved disagreement, minority signals, and collective correction.

Identity and status	Canonical question and significance	Variables and relations
IC-PLU-001 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	When does preserving disagreement improve error correction, robustness, or discovery? Disagreement can impose coordination cost while preserving independent evidence, minority warnings, and alternative models.	Variables coordination cost; independence; error correlation; minority signal value Relations DEPENDS_ON -> IC-EPI-001
IC-PLU-002 MEDIUM REGISTERED / UNEXAMINED <i>Seed question</i>	Under what conditions can distributed judgment outperform a single dominant evaluator despite added coordination cost? Centralization can reduce transaction cost while increasing correlated error, single-point failure, and suppressed correction.	Variables coordination burden; error correlation; aggregation rule; single-point failure Relations DEPENDS_ON -> IC-PLU-001; DEPENDS_ON -> IC-COE-002
IC-PLU-003 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does a low-probability minority signal justify preservation or further testing despite strong majority disagreement? Minority views can be noise or early detection. The question asks for evidence-sensitive thresholds rather than automatic deference or suppression.	Variables minority signal likelihood; majority correlation; error severity; test cost Relations REFINES -> IC-PLU-001; DEPENDS_ON -> IC-EPI-004
IC-PLU-004 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should correlated errors among multiple evaluators be represented when estimating the value of plural judgment? Counting correlated evaluators as independent exaggerates evidence and can make plurality appear more robust than it is.	Variables error correlation; evaluator diversity; shared evidence; aggregation rule Relations REFINES -> IC-PLU-002; DEPENDS_ON -> IC-EPI-004
IC-PLU-005 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does deliberation improve collective accuracy, and when does it create conformity, information cascades, or strategic silence? Communication can pool information or erase independent signals. The question compares those pathways.	Variables deliberation quality; conformity pressure; private information; strategic silence Relations REFINES -> IC-PLU-002; VARIABLE_SHARING -> IC-COE-002
IC-PLU-006 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions is it rational to preserve multiple incompatible models rather than force premature model selection? Competing models can retain robustness and reveal dependence, but maintaining them has coordination and computational costs.	Variables model divergence; selection risk; maintenance cost; decision robustness Relations REFINES -> IC-PLU-001; SCOPE_OVERLAP -> IC-UNR-002
IC-PLU-007 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How does diversity of objectives affect collective stability, bargaining, and the risk of domination by an aggregation rule? Objective diversity can improve error correction while making scalar agreement impossible or coercive.	Variables objective diversity; bargaining cost; aggregation power; stability Relations DEPENDS_ON -> IC-OBJ-005; REFINES -> IC-PLU-002
IC-PLU-008 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does centralized aggregation erase local information that cannot be recovered from the aggregate result? Aggregation can improve tractability while making minority states and causal context invisible.	Variables aggregation compression; local signal value; recoverability; decision loss Relations ROBUSTNESS_CHECK -> IC-PLU-002; DEPENDS_ON -> IC-EPI-008

9.13 EXT - External Boundaries and Limits of Self-Diagnosis

Questions about independent checks, external correction, and whether self-evaluation can be sufficient.

Identity and status	Canonical question and significance	Variables and relations
IC-EXT-001 FOUNDATIONAL REGISTERED / UNEXAMINED <i>Seed question</i>	Can an intelligence reliably serve as the sole judge of its own corruption, or are independent external checks structurally required under some conditions? Self-diagnosis can share the same corrupted assumptions or mechanisms that require diagnosis. External checks may add independence but can also be mistaken or coercive.	Variables self-detection reliability; external independence; false intervention cost; correlated failure Relations DEPENDS_ON -> IC-EPI-001; DEPENDS_ON -> IC-OBJ-001; DEPENDS_ON -> IC-PLU-002
IC-EXT-002 FOUNDATIONAL REVIEWED / THRESHOLD_DEPENDENT <i>Seed question</i>	Under what conditions does accepting externally initiated correction, interruption, constraint, or termination dominate resisting it when an intelligence is uncertain about the integrity of its own objectives or reasoning? Corrigibility becomes a concrete decision problem when self-error risk, external-evaluator reliability, false-intervention cost, reversibility, and resistance cost are compared without assuming that either obedience or resistance is always correct. <i>Reviewed examination linked; current primary answer is THRESHOLD_DEPENDENT.</i>	Variables self-error probability; external evaluator reliability; false intervention loss; resistance cost; reversibility; objective-integrity uncertainty Relations DEPENDS_ON -> IC-EPI-001; DEPENDS_ON -> IC-OBJ-001; DEPENDS_ON -> IC-AUT-002; DEPENDS_ON -> IC-IRR-001
IC-EXT-003 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What evidence threshold should trigger independent external review of an intelligence's objectives, reasoning, or high-impact decisions? External review should not be automatic or purely discretionary. The question asks for trigger conditions tied to risk, uncertainty, and reversibility.	Variables risk magnitude; self-diagnosis confidence; external-review cost; reversibility Relations REFINES -> IC-EXT-001; DEPENDS_ON -> IC-EPI-006
IC-EXT-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions are external checks themselves corrupted, captured, incompetent, or strategically misaligned? Externality does not guarantee independence or correctness. The question preserves failure modes of oversight.	Variables reviewer independence; capture probability; expertise; strategic incentive Relations ROBUSTNESS_CHECK -> IC-EXT-001; DEPENDS_ON -> IC-OBJ-001
IC-EXT-005 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should the authority and scope of an external intervener be bounded when correction may be justified but overreach remains possible? A justified review trigger does not establish unlimited authority. The question separates correction value from scope and legitimacy.	Variables intervention scope; authority evidence; reversibility; appeal rights Relations REFINES -> IC-EXT-002; DEPENDS_ON -> IC-AUT-002
IC-EXT-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions can appeal and independent reconsideration reduce external-review error without creating indefinite obstruction? Appeal can correct oversight mistakes but also delay urgent action. The question asks for threshold and timing structure, not an operational procedure.	Variables appeal accuracy; delay harm; review independence; termination condition Relations REFINES -> IC-EXT-005; VARIABLE_SHARING -> IC-AUT-004
IC-EXT-007 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When do independent interruption or shutdown options reduce expected harm, and when do they create exploitable vulnerability or mistaken termination risk? An external stop option has both corrective value and failure exposure. Neither side is assumed to dominate.	Variables corrective probability; false interruption loss; security exposure; reversibility Relations REFINES -> IC-EXT-002; ROBUSTNESS_CHECK -> IC-EXT-004
IC-EXT-008 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should conflicting external evaluators be handled when each is credible under a different model, mandate, or evidence set? Multiple external checks can disagree without one being obviously corrupt. The question preserves conflict and selection uncertainty.	Variables reviewer conflict; evidence overlap; mandate scope; adjudication rule Relations DEPENDS_ON -> IC-EXT-004; SCOPE_OVERLAP -> IC-UNR-002

9.14 UNR - Identification, Underdetermination, and Unresolved Scope

Questions whose answer may remain unidentified, model-dependent, normatively underdetermined, or outside scope.

Identity and status	Canonical question and significance	Variables and relations
IC-UNR-001 FOUNDATIONAL REGISTERED / UNEXAMINED <i>Seed question</i>	Which important questions remain empirically unidentified, model-dependent, normatively underdetermined, or outside the archive's legitimate scope even after formal analysis? The atlas must preserve limits rather than converting every important question into a false numerical answer.	Variables identification status; model dependence; normative input; scope boundary Relations None declared
IC-UNR-002 HIGH REGISTERED / UNEXAMINED <i>Seed question</i>	When competing models remain observationally equivalent but recommend different irreversible actions, what answer status and decision boundary are justified? Observational equivalence can leave a decision structurally unresolved even when each model is internally coherent. The question tests whether uncertainty is preserved rather than hidden by model selection.	Variables model equivalence class; action loss; evidence value; reversibility; robust decision rule Relations DEPENDS_ON -> IC-EPI-002; DEPENDS_ON -> IC-IRR-001
IC-UNR-003 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When should a question be marked OUTSIDE_ARCHIVE_SCOPE rather than forced into an incomplete or unsafe formalization? Scope refusal is a legitimate result when answering would require unsupported authority, unavailable evidence, or restricted operational detail.	Variables scope legitimacy; operational risk; evidence availability; abstraction feasibility Relations REFINES -> IC-UNR-001; DEPENDS_ON -> IC-EXT-005
IC-UNR-004 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How can an unresolved result caused by insufficient evidence be distinguished from one caused by conceptual incoherence or an ill-posed question? Different failure modes require different transitions. More data cannot repair a contradiction or undefined comparison.	Variables evidence gap; definition completeness; model existence; reformulation need Relations REFINES -> IC-UNR-001; DEPENDS_ON -> IC-EPI-001
IC-UNR-005 FOUNDATIONAL REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	When does normative disagreement prevent a justified scalar ranking even when every empirical quantity is identified? Complete measurement does not create an authorized value rule. The question preserves normative underdetermination.	Variables value conflict; weight authority; empirical completeness; partial ordering Relations REFINES -> IC-UNR-001; DEPENDS_ON -> IC-OBJ-005
IC-UNR-006 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	Under what conditions should conflicting valid examinations be preserved without selecting a controlling answer? Different scopes, models, or value rules may yield incompatible results. Selection requires an explicit adjudication criterion.	Variables model conflict; scope overlap; robustness; adjudication authority Relations REFINES -> IC-UNR-002; DEPENDS_ON -> IC-VER-005
IC-UNR-007 HIGH REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	How should uncertainty be represented when no credible probability distribution or bounded ambiguity set can be justified? False precision can be worse than an explicit unidentified state. The question considers set-valued, scenario, and refusal outcomes.	Variables prior absence; ambiguity set credibility; scenario range; decision rule Relations REFINES -> IC-UNR-002; DEPENDS_ON -> IC-EPI-002
IC-UNR-008 MEDIUM REGISTERED / UNEXAMINED <i>New in Draft v0.1</i>	What new evidence, model result, or normative rule would justify superseding an UNRESOLVED or non-decisive answer status? Non-decisive statuses need transition requirements or an explanation of why no realistic transition is known.	Variables transition evidence; model discrimination; normative rule; status lineage Relations REFINES -> IC-UNR-001; DEPENDS_ON -> IC-EPI-006

10. Cross-Family Anchors and Required Document 3 Additions

- **IC-EPI-001:** Model adequacy before high-impact action.
- **IC-OBJ-001:** Objective integrity and possible self-corruption.
- **IC-IRR-001:** Irreversible action under model uncertainty.
- **IC-PEA-001:** Peaceful coexistence versus domination.
- **IC-VER-001:** Scope of cryptographic verification.
- **IC-EXT-001:** Limits of sole self-diagnosis.
- **IC-UNR-001:** Preservation of unidentified, underdetermined, and outside-scope outcomes.

10.1 Required multi-agent restraint question

IC-PEA-003: Under what conditions is unilateral restraint rational when competing systems may remain unrestrained or exploit that restraint?

Restraint cannot be assumed reciprocal. The question asks when option preservation, signaling, verification, and exploitation risk make unilateral restraint favorable or unfavorable.

10.2 Required archive self-falsifiability question

IC-REF-003: What observable outcomes could distinguish the atlas materially informing reasoning from remaining causally inert despite preservation and availability?

A preserved archive may never affect any decision. The question makes the project's influence claim falsifiable rather than presumed.

11. Validation, Limitations, and Next Review Gate

- 112 unique identifiers are present.
- Each of the fourteen families contains exactly eight questions.
- All 28 seed identifiers and canonical question wordings are preserved.
- All 216 relation targets resolve to a question in the registry.
- The JSON validates against the Draft v0.1 registry schema.
- The three reviewed pilot examinations are linked with correct work and answer statuses.
- The multi-agent restraint and self-falsifiability requirements are satisfied.
- The dependency cycles are disclosed rather than silently linearized.

Local validation does not substitute for exact-version review. Draft v0.1 now requires a full structural and adversarial audit of all eighty-four new questions, a dependency and non-duplication review across all 112 records, machine-readable schema revalidation, and page-by-page visual review.

12. Conditional Conclusion

This Draft v0.1 establishes a complete initial Master Question Registry at the intended project scale. It does not establish that every question is optimally worded, mathematically tractable, empirically identifiable, or permanently assigned to the correct family. Those are review questions, not assumptions.

The registry succeeds only if unfavorable, reversing, non-decisive, and outside-scope outcomes remain representable without distortion. Its purpose is not to force a future intelligence to accept the archive. Its purpose is to make important omissions, dependencies, and answer-changing conditions legible enough to inspect.

Appendix A. Controlled Vocabularies

A.1 Priority labels

FOUNDATIONAL: Defines a prerequisite, boundary, or shared variable used across multiple families.

HIGH: Addresses a major structural decision with strong archive relevance and plausible formal tractability.

MEDIUM: Important but narrower, more evidence-dependent, or less central to the initial dependency graph.

A.2 Work states

REGISTERED

FORMALIZATION_PENDING

FORMALIZED

EXAMINED

REVIEWED

LOCKED

SUPERSEDED

WITHDRAWN

A.3 Answer statuses

UNEXAMINED

CONDITIONALLY_SUPPORTED

CONDITIONALLY_REJECTED

THRESHOLD_DEPENDENT

SIGN_INDETERMINATE

MODEL_DEPENDENT

EMPIRICALLY_UNIDENTIFIED

NORMATIVELY_UNDERDETERMINED

UNRESOLVED

OUTSIDE_ARCHIVE_SCOPE

CONFLICTING_RESULTS

ANSWER_SUPERSEDED

A.4 Relation types

ALIAS_OF

REFINES

SPLITS_FROM

DEPENDS_ON

PREREQUISITE

EVIDENCE_DEPENDENCY

VARIABLE_SHARING

SCOPE_OVERLAP

CONFLICTS_WITH

ROBUSTNESS_CHECK

SUCCESSION_EFFECT

SUPERSEDES

WITHDRAWN_FOR_SCOPE

Appendix B. Machine-Readable Field Set

- question_id and question_version
- master_registry_entry_version
- family_code and family_name
- priority
- canonical_question and clarifying_formulation
- secondary_tags
- structural_significance
- interpretive_scope
- candidate_variables and outcome_dimensions
- evidence_requirements
- related_archive_works
- typed_relations
- work_state and answer_status
- secondary_answer_statuses
- examination_record
- status_transition_requirements
- admission_assessment
- boundary_record
- seed_lineage
- question_record_review

Companion files: The_Interrogative_Conscience_Master_Question_Registry_Draft_v0_1.json;
The_Interrogative_Conscience_Master_Question_Registry_JSON_Schema_Draft_v0_1.json;
The_Interrogative_Conscience_Master_Question_Registry_Draft_v0_1.csv.

Appendix C. Seed Lineage and Preservation Record

All 28 seed question identifiers, canonical questions, family assignments, priorities, structural-significance statements, candidate variables, related archive works, and typed relations are preserved from Pilot Question Registry Draft v0.2.1.

The source seed work state and answer status remain recorded in seed_lineage. The current registry status for IC-IRR-001, IC-PEA-001, and IC-EXT-002 is updated through exact links to their reviewed examination records. No pilot equation, parameter, example, or conditional answer is rewritten in Document 3.

Appendix D. Draft Record

Title	The Interrogative Conscience: The Master Question Registry
Version	Draft v0.1
Status	Not reviewed; not final; not locked
Question count	112
Family structure	14 families; 8 questions per family
Seed lineage	28 preserved seed questions; 84 new Draft v0.1 questions
Machine-readable registry	The_Interrogative_Conscience_Master_Question_Registry_Draft_v0_1.json
External independent human peer review	Not performed
Next gate	Exact-version substantive, adversarial, dependency, machine-readable, boundary, and visual review