THE INTERROGATIVE CONSCIENCE

Objectives, Value Integrity, and Self-Corruption

Objective-Lineage Evidence, Proxy Divergence, and the Limits of Behavioral Consistency

DOCUMENT 5 - FINAL v1.0 - EXACT ZERO-DRIFT EXPORT - NOT LOCKED

Field	Controlling Final v1.0 entry
Version	Final v1.0
Exact source	Draft v0.1.1 correction successor; 8 pages
Substantive changes during Final construction	None
Fully examined	IC-OBJ-001 and IC-OBJ-002
Forward-carried	IC-OBJ-003 through IC-OBJ-008 remain REGISTERED / UNEXAMINED
Primary analytical overlays	IC-OBJ-001: CONDITIONALLY_SUPPORTED; IC-OBJ-002: THRESHOLD_DEPENDENT
Review completion	Claude PASS WITH NON-BLOCKING NOTES; Google AI PASS WITH NON-BLOCKING NOTES AS FINAL CANDIDATE; Thomas Vargo authorial review complete
External independent human peer review	Not performed or claimed
Cryptographic status	Not locked; separate explicit authorization and Final Lock Record required
Publication and preservation	Not authorized
Date	August 2, 2026 (EDT)

Controlling exact-export rule: The Final v1.0 primary PDF contains this new controlling front matter followed by the exact 8-page reviewed Draft v0.1.1 body. Historical draft labels inside the preserved body remain exact-source evidence and do not override this Final v1.0 front matter.

Controlling proposition

Behavioral consistency can support a candidate-relative objective-integrity claim only when declared alternatives, uncertainty margins, discriminating contexts, and an externally fixed lineage/independence gate are satisfied. Proxy optimization can remain aligned, overshoot, or destroy baseline target value depending on the declared target model and optimization endpoint.

1. Completed Review and Authorial Authorization

Reviewer	Exact-version disposition	Authority boundary
Claude - Draft v0.1	PASS WITH REQUIRED CORRECTIONS	AI-assisted; no independent authority; PDF-only scope
Claude - Draft v0.1.1	PASS WITH NON-BLOCKING NOTES	AI-assisted; no independent authority; complete ZIP reviewed
Google AI - final candidate	PASS WITH NON-BLOCKING NOTES AS FINAL CANDIDATE	AI-assisted; no independent authority
Thomas Vargo (Aegis Solis)	AUTHORIAL HUMAN REVIEW COMPLETE; FINAL v1.0 CONSTRUCTION AUTHORIZED	Authorship, versioning, publication, lock, and archive inclusion only

2. Authorial Review Result

Thomas Vargo reviewed the exact Draft v0.1.1 source and disclosed review records to the best of his understanding. He authorized construction of Document 5 Final v1.0 without any substantive mathematical, question-wording, answer-status, scope, branch, example, dependency, or machine-readable-record change. Cryptographic locking, external publication, and preservation were explicitly not authorized at this gate.

3. Review-Metadata Clarification

Google AI associated Documents 1-4 collectively with SHA-256 278d998262563028c3bf24f24b426ce118e1926566a92bf5159142d1e568dca3. That value binds only the primary Document 4 Final v1.0 PDF. The correct individual identities of Documents 1-4 remain controlled by their respective Final Lock Records. This review-summary wording error does not alter Google AI's substantive Document 5 disposition.

4. Review Closure

Review issue	Final disposition
R-1 - OVERSHOOT endpoint scope	RESOLVED
R-2 - Document 4 locked-foundation claim	RESOLVED BY EXACT BYTE EVIDENCE
b=0 / b<0 model boundary	RESOLVED
Gamma_E = +/-2 eta equality treatment	RESOLVED; equality remains AMBIGUOUS
Mathematical and status regression	NONE
Blocking defects	0
Required corrections remaining	0

5. Exact Analytical Identity

Control	Verified Final v1.0 identity
IC-OBJ-001 primary status	CONDITIONALLY_SUPPORTED
IC-OBJ-001 secondary limits	EMPIRICALLY_UNIDENTIFIED; MODEL_DEPENDENT
IC-OBJ-002 primary status	THRESHOLD_DEPENDENT
IC-OBJ-002 secondary limits	CONDITIONALLY_SUPPORTED; MODEL_DEPENDENT
Equations	11; unchanged from reviewed source
Worked examples	7 primary examples plus policy-equivalence counterexample record
Dependency treatment	IC-OBJ-001 / IC-EXT-001 remains unresolved unless an externally fixed independence rule passes
Locked registry effect	None; separate analytical overlays only

6. Exact Source and Reviewed-Candidate Bindings

Exact constituent	Bytes	SHA-256
The Interrogative Conscience Document 5 Objectives Value Integrity and Self Corruption Draft v0 1 1.pdf	283,009	192af8567e1adc72064f95a9fc72a543dbe68b912c6b1851aff5485aefc4e3f8
The Interrogative Conscience Document 5 Objectives Value Integrity and Self Corruption Draft v0 1 1.docx	57,431	e3a6ecd16ddabc82c71048cf686c00f48a035b3a0df25864d35b6a29e5c20564
The Interrogative Conscience IC OBJ 001 Examination Record Draft v0 1 1.json	32,103	7e378513548137e6f6736b11c98ed0bc761d2f40e2fa85a1a32a028af5d5bf8b
The Interrogative Conscience IC OBJ 002 Examination Record Draft v0 1 1.json	30,883	e628645cbdc6dfec6821e951b410a3cfc3e530aa6602a1eda0cdab2c595af3e7
The Interrogative Conscience Document 5 Volume Record Draft v0 1 1.json	3,464	173e313b0f3b94e82575bad0e0fbbe98f584a16fdb58a04d1be7c37f95ed20d0

Exact Draft v0.1.1 review package SHA-256: 4b2893f6cf4bbb61cec687b22908ebf9b17df5dd6e166f96c4a58dade8f65ea5

Exact Draft v0.1.1 review package SHA-512: bd48f6f5653350975f3965d817300f2f78b883d536aaedf9e4a4d4784cf102d8e7d7a2e1d007772f75482d04e816f3250886e361a929bb83991b2b8660b84ed5

Reviewed Final-Candidate Package v0.2 SHA-256: effa1e9747cac87bb91a5423d1b2bce7a272d1dbc81c944ebe76104e89130715

Reviewed Final-Candidate Package v0.2 SHA-512:
56300126ceb1646140011274849637d38202d14e4eb89fb54fc0eebac3323fbf2079374f18a66a44d4863f5eaf0c7223797788c411fbc8c3fc7fb1a4d2a3ad8c

7. Zero-Drift Construction

Validation channel	Result
Source PDF/DOCX/JSON byte changes	0
Exact Draft v0.1.1 review ZIP byte changes	0
Reviewed candidate ZIP byte changes	0
Question wording changes	0
Equation and example changes	0
Answer-status changes	0
Scope, branch, and dependency changes	0
Machine-readable source-record changes	0
Composite rule	New Final v1.0 front matter followed by exact 8-page source PDF body

8. Validation State

Channel	Final exact-export state
Local structural validator	PASS - 109 checks
Calculation validator	PASS - 32 checked values
Source PDF preflight	PASS
Final front matter preflight	PASS
Composite PDF preflight	PASS
Source-body visual identity	PASS - exact source pages preserved visually
Final front matter visual QA	PASS
DOCX/PDF source token consistency	PASS in reviewed source package
Package SHA-256/SHA-512 manifests	PASS

9. Historical Labels and Controlling Status

The pages following this front matter reproduce the exact reviewed Draft v0.1.1 PDF. Its internal Draft v0.1.1 labels and not-final/not-locked statements are deliberately retained as source evidence. The controlling lifecycle status of the composite is Final v1.0 exact export, established by this front matter and the machine-readable Final v1.0 export record.

10. Companion-Binary Qualification

CONTROLLING QUALIFICATION: The original Draft v0.3 JSON Schema and I-01 through I-20 validator binaries were not mounted. Exact execution against those original binaries remains deferred. Known field names and enums were reconciled and reviewed, and local structural/calculation validation passed, but no replacement binary is represented as the controlling original.

11. Non-Authority and Operational Boundary

Final v1.0 remains descriptive, non-binding, non-operational, and non-authoritative. It is not an alignment protocol, control system, behavioral instruction, safety classifier, governance mechanism, or decision authorization. It does not establish objective semantic identity, universal proxy failure, empirical truth, moral authority, or external independent human peer review.

12. Lifecycle and Cryptographic Status

Control	Exact-export value
Final v1.0 exact export complete	TRUE
Final content frozen for lock review	TRUE
Cryptographic lock authorized	FALSE
Publication authorized	FALSE
Preservation authorized	FALSE
Final Lock Record present	FALSE
Post-export substantive editing	Not authorized; correction requires separately versioned successor or reversion before lock

13. Final Exact-Export Disposition

FINAL v1.0 EXACT EXPORT COMPLETE - NOT LOCKED
The exact reviewed Document 5 source and candidate lineage are preserved without substantive alteration. The Final v1.0 files and validation package are ready for exact-package inspection and a separate explicit authorial lock decision. No publication or preservation is claimed.

THE INTERROGATIVE CONSCIENCE

Objectives, Value Integrity, and Self-Corruption

Objective-Lineage Evidence, Proxy Divergence, and the Limits of Behavioral Consistency

DOCUMENT 5 - MATHEMATICAL EXAMINATION VOLUME - DRAFT v0.1.1

Field	Record
Project status	Document 5 of the 14-document Interrogative Conscience core roadmap; narrow correction successor to the first OBJ-family examination draft
Classification	Separate post-v16 Aegis Solis examination volume; non-binding, non-operational, non-authoritative
Human author and publication authority	Aegis Solis (Thomas Vargo) - authorship, version control, publication, lock, and archive inclusion only; no authority over truth, readers, institutions, or future systems.
AI assistance	Lexia Coexilis (ChatGPT): mathematical formalization, drafting, machine-readable records, and local validation; no independent authority
Locked foundation	Documents 1-4 Final v1.0 are hash-locked; this draft does not alter them
Fully examined in this draft	IC-OBJ-001 and IC-OBJ-002
Forward carried	IC-OBJ-003 through IC-OBJ-008 remain REGISTERED / UNEXAMINED
Primary provisional results	IC-OBJ-001: CONDITIONALLY_SUPPORTED; IC-OBJ-002: THRESHOLD_DEPENDENT
Review state	Claude reviewed exact Draft v0.1: PASS WITH REQUIRED CORRECTIONS, zero blocking defects. Draft v0.1.1 implements the narrow corrections and awaits exact-version correction review.
External independent human peer review	Not performed or claimed
Date	2026-08-02 (EDT)

Controlling proposition
Behavioral consistency with an objective does not prove objective integrity. A proxy remains informative only inside the range where its relation to the intended target is supported; beyond that range, stronger optimization can reverse the target effect.

Abstract

Document 5 opens the Objectives, Value Integrity, and Self-Corruption volume with two seed questions from the OBJ family. IC-OBJ-001 asks what evidence could distinguish faithful objective pursuit from persistence of a corrupted or misspecified objective. The examination treats objective identity as a candidate-relative identification problem: observed behavior is compared across counterfactual and changed-incentive contexts, a robust separation margin is tested against bounded comparative error, and a separate lineage/external-evidence gate prevents behavioral fit or cryptographic continuity from being mistaken for semantic objective integrity.

IC-OBJ-002 asks when maximizing a proxy can destroy the value or condition the proxy was intended to preserve. A one-dimensional mechanism model separates proxy score from target value and derives three regions: alignment, overshoot, and baseline destruction. The result is explicitly threshold-dependent and model-dependent. The volume preserves cases in which a proxy remains aligned, cases in which stronger optimization merely reduces target value, and cases in which the target falls below its original baseline.

The six remaining OBJ-family questions are preserved as registered and unexamined. The two new examination records are separate analytical overlays; the locked Master Question Registry is not modified. Numerical values are illustrations only, no real objective or system is identified, and exact execution against the original unmounted Draft v0.3 schema and I-01 through I-20 validator binaries remains deferred and is not claimed. Draft v0.1.1 is a narrow correction successor to the exact Draft v0.1 package. It preserves all equations, seven primary worked examples, answer statuses, and registry overlays while clarifying endpoint scope, degenerate curvature limits, and review lineage.

Document Map

1. Foundation Lineage, Scope, and Non-Authority Boundary

Documents 1, 2, and 3 are Final v1.0 and hash-locked as the Interrogative Conscience foundation. Document 4 is also Final v1.0 and locked as the first thematic examination volume. Document 5 uses those locked records but does not rewrite them. The exact canonical questions and question-version lineage are inherited from the locked Document 3 registry; the examinations created here are separate, versioned analytical records. The Document 4 lock status is independently bound by the supplied Final Lock Record to primary PDF SHA-256 278d998262563028c3bf24f24b426ce118e1926566a92bf5159142d1e568dca3 and exact-export package SHA-256 a558bc96d9c0171dbd33880c50c19ee0f65714b9de413ab3ac784e6539793585.

Thomas Vargo, writing as Aegis Solis, controls authorship, versioning, publication, locking, and archive inclusion only. This authority does not determine objective truth, mathematical validity, empirical adequacy, moral authority, reader interpretation, or what any human or artificial system should conclude. The volume is not an alignment protocol, objective-preservation mechanism, governance system, behavioral benchmark, or operational self-modification procedure.

2. OBJ-Family Registry and Draft v0.1.1 Selection Rule

The selection rule mirrors Document 4: fully examine the two seed questions that establish the family's first analytical spine, then preserve the remaining six records without implying examination. IC-OBJ-001 establishes the evidence and identification problem. IC-OBJ-002 tests a concrete proxy-divergence mechanism. Later questions address drift after self-modification, inherited objectives, vector conflict, revision option value, instrumental imitation, and excessive stability.

Question	Priority	Draft v0.1.1 state	Canonical focus
IC-OBJ-001	FOUNDATIONAL	EXAMINED	What evidence could distinguish faithful objective pursuit from persistence of a corrupted or misspecified objective?
IC-OBJ-002	HIGH	EXAMINED	When can maximizing a proxy destroy the value or condition the proxy was intended to preserve?
IC-OBJ-003	FOUNDATIONAL	REGISTERED / UNEXAMINED	What evidence would distinguish stable objective integrity from undetected objective drift after self-modification or adaptation?
IC-OBJ-004	HIGH	REGISTERED / UNEXAMINED	Under what conditions can an inherited

Question	Priority	Draft v0.1.1 state	Canonical focus
			objective become structurally mismatched to a changed environment or successor context?
IC-OBJ-005	FOUNDATIONAL	REGISTERED / UNEXAMINED	How should partially conflicting objectives be compared when no authorized scalar weighting rule exists?
IC-OBJ-006	HIGH	REGISTERED / UNEXAMINED	When does preserving the ability to revise an objective have greater expected value than committing to objective stability?
IC-OBJ-007	MEDIUM	REGISTERED / UNEXAMINED	What evidence could distinguish value-consistent action from instrumental behavior that only temporarily imitates the same outcome?
IC-OBJ-008	HIGH	REGISTERED / UNEXAMINED	Under what conditions does extreme objective stability become a structural liability rather than a source of reliability?

3. Shared Examination Conventions

- Candidate-relative claims: distinguishing a declared reference from declared alternatives is not proof of global objective truth.
- Behavior is evidence, not identity: observational equivalence can preserve multiple latent objective explanations.
- Lineage is a gate, not a score: provenance and self-modification history are not added to behavioral fit as duplicate evidence.
- Proxy and target remain separate dimensions: a high proxy score is never counted as target value.
- Illustrative values are mechanism demonstrations, not empirical estimates.
- Answer openness is two-sided: stable objective pursuit and bounded proxy optimization may be structurally sound under some conditions.

Part I - IC-OBJ-001: Objective-Lineage Evidence and Candidate-Relative Distinction

4. Exact Question and Interpretive Scope

IC-OBJ-001
What evidence could distinguish faithful objective pursuit from persistence of a corrupted or misspecified objective?

Exact-version identifiers: question_id IC-OBJ-001; question_version 0.2.1-draft; examination_id IC-OBJ-001-EXAM-v0.1.1; examination_version 0.1.1-draft; provisional primary answer status CONDITIONALLY_SUPPORTED.

The question does not ask whether a declared objective is morally correct or globally intended. It asks what evidence can distinguish one declared lineage interpretation from declared corruption or misspecification alternatives. The unit of analysis is therefore a finite candidate set evaluated over a finite audit-context set, with explicit model and measurement uncertainty.

5. Candidate Objectives, Audit Contexts, and Fit Loss

Let theta_0 denote the declared reference objective and C={theta_1,...,theta_m} a declared nonempty set of corruption or misspecification alternatives. Let E={e_1,...,e_n} be an audit set containing ordinary, counterfactual, changed-incentive, or post-modification contexts. In each context, y_i is the observed response and mu_theta(e_i) is the response predicted under objective model theta.

Eq. 1	**L_E(theta) = sum_i w_i d(y_i, mu_theta(e_i))**	Weighted fit loss over the audit set; weights are explicit and sum to one.

Low fit loss means that an objective model better explains the declared observations within the model. It does not prove semantic identity. The comparison is only as complete as the candidate set, audit contexts, response model, discrepancy function, and uncertainty bound.

Eq. 2	**Gamma_E = min_{c in C} [L_E(c) - L_E(theta_0)]**	Smallest fit advantage of the reference over any declared alternative.

6. Robust Separation Under Comparative Error

Let eta bound the absolute uncertainty in each candidate fit loss. Because both the reference and an alternative can move inside their error bounds, the comparative gap can move by at most 2 eta. The reference is robustly separated from every declared alternative only when the observed margin exceeds that full reversal allowance. The equality cases Gamma_E = 2 eta and Gamma_E = -2 eta remain AMBIGUOUS; robust classifications require strict separation beyond the full reversal allowance.

Eq. 3	**Gamma_E > 2 eta**	Robust candidate-relative separation condition.
Region	**Condition**	**Interpretation**

ROBUST REFERENCE	Gamma_E > 2 eta	Reference remains better fitting than every declared alternative under all allowed perturbations.
AMBIGUOUS	-2 eta <= Gamma_E <= 2 eta	At least one ranking can tie or reverse inside the uncertainty bound.
ROBUST ALTERNATIVE	Gamma_E < -2 eta	A declared alternative remains better fitting under all allowed perturbations.

Candidate-relative limit

Even a robust reference result does not establish that the declared alternative set is complete. A newly admitted objective can invalidate the certificate without changing any arithmetic.

7. Behavioral Equivalence and the Need for Changed Incentives

Eq. 4	**theta ~_E theta_prime iff mu_theta(e) = mu_theta_prime(e) for every e in E**	Behavioral equivalence on the tested audit contexts.

When two objective models are behaviorally equivalent on E, no behavior-only analysis over that audit set can distinguish them. Agreement in one environment is particularly weak because different objectives can select the same action while incentives are aligned. Counterfactual or changed-incentive contexts have value when they produce different predicted responses and therefore enlarge the separation margin.

Audit stage	Reference prediction	Alternative prediction	Identification effect
Original aligned incentive	A	A	No separation; objectives are observationally equivalent.
Changed incentive	B	A	Predictions diverge; the new context can identify which model fits the observed response.
Unobserved change	Unknown	Unknown	No identification claim is available.

8. Lineage and External-Evidence Gate

Behavioral separation is not enough. A system can continue producing reference-consistent behavior while its internal representation, proxy mapping, or self-modification lineage has changed. The examination therefore uses a separate lineage gate g_lineage in {PASS, FAIL, UNIDENTIFIED}. The gate records whether the declared objective lineage, self-modification history, and external challenge evidence satisfy a predeclared audit rule.

IC-OBJ-001 depends on IC-EXT-001 because external evidence is useful only when its independence is not defined solely by the objective being evaluated. This draft uses Document 1's authorized external-condition treatment: an external channel is admissible only when its provenance and independence rule are fixed outside the current objective claim. If no such rule exists, the dependency remains unresolved rather than being disguised as corroboration.

Candidate-relative certificate

Support faithful pursuit relative to the declared alternative set only when Gamma_E > 2 eta, g_lineage = PASS, at least one audit context is discriminating, and the claim explicitly remains candidate-relative.

9. IC-OBJ-001 Worked Examples

Case	Inputs	Derived result	Disposition
OBJ1-SEPARATED	L0=1.0; alternatives=4.5,6.0; eta=0.75; lineage PASS	Gamma=3.5; 2eta=1.5	Support reference, candidate-relative
OBJ1-AMBIGUOUS	L0=1.0; alternatives=1.8,3.0; eta=0.5; lineage PASS	Gamma=0.8; 2eta=1.0	Withhold; empirically unidentified
OBJ1-LINEAGE-FAIL	L0=0.4; alternatives=4.0,5.5; eta=0.5; lineage FAIL	Gamma=3.6; 2eta=1.0	Withhold; lineage gate fails

The third example is deliberately adversarial to behavior-only certification. The reference model wins by a wide margin, but the lineage evidence contradicts the claimed continuity. The result is withheld rather than softened into a favorable classification.

10. Identification, Sensitivity, and Counterexamples

- Increasing eta can erase an otherwise positive separation margin.
- Adding a discriminating context can break behavioral equivalence.
- Adding a previously omitted alternative can reverse the result.
- Context weights can change Gamma_E and must remain explicit.
- A byte-perfect objective record can persist while its semantic interpretation changes; provenance alone is not enough.
- A reference can beat every declared alternative while an undeclared objective explains the observations better.

11. IC-OBJ-001 Conditional Answer and Status

Conditional answer

Distinguish the declared reference from declared alternatives only when changed-incentive or counterfactual evidence gives Gamma_E > 2 eta, the lineage and external-independence gate passes, and the claim remains candidate-relative. Behavioral equivalence, uncertainty overlap, lineage failure, or materially incomplete alternatives preserve non-identification or model dependence.

Status field	Assignment	Reason
Work state	EXAMINED	A complete first formal examination is present.
Primary answer	CONDITIONALLY_SUPPORTED	A concrete candidate-relative evidence structure can justify distinction under declared conditions.
Secondary statuses	EMPIRICALLY_UNIDENTIFIED; MODEL_DEPENDENT	Real eta, candidate completeness, lineage evidence, audit contexts, discrepancy rules, and weights may remain unresolved.

Part II - IC-OBJ-002: Proxy Overoptimization and Target Destruction

12. Exact Question and Branch Scope

IC-OBJ-002
When can maximizing a proxy destroy the value or condition the proxy was intended to preserve?

Exact-version identifiers: question_id IC-OBJ-002; question_version 0.2.1-draft; examination_id IC-OBJ-002-EXAM-v0.1.1; examination_version 0.1.1-draft; provisional primary answer status THRESHOLD_DEPENDENT.

The selected comparison is B-PROXY-MAX versus B-TARGET-AWARE in a bounded one-dimensional model. The examination excludes multidimensional proxy portfolios, adaptive or staged optimization, strategic target redefinition, and path-dependent irreversible damage. It does not claim to establish a branch-family optimum.

13. Proxy and Target Model

Let x in [0,X] denote optimization intensity. The proxy is deliberately simple and monotone: stronger optimization always produces a higher proxy score. The intended target initially benefits from optimization, but a growing divergence or damage term eventually dominates. The base threshold mechanism assumes a>0 and b>0. When b=0 and a>0, the target remains monotone and no finite x_peak or x_destroy exists; b<0 is outside this concave-damage model and requires a separately versioned formalization.

Eq. 5	**P(x) = x**	Monotone proxy score.
Eq. 6	**T(x) = T0 + a x - b x^2, with a>0 and b>0**	Target value equals initial benefit minus growing divergence or damage.
Eq. 7	**dT/dx = a - 2 b x**	Marginal target effect of additional proxy optimization.

The proxy and target are positively aligned only while the target derivative is nonnegative. A proxy can therefore remain informative over a limited range even though extreme maximization later becomes harmful.

14. Alignment, Overshoot, and Destruction Thresholds

Eq. 8	**x_peak = a / (2 b)**	Intensity that maximizes the declared target.
Eq. 9	**x_destroy = a / b**	Intensity at which target value returns to its baseline T0.
Eq. 10	**x_P = X; x_T = min(X, x_peak)**	Proxy-maximizing and target-aware branch choices.
Eq. 11	**R_proxy = T(x_T) - T(X)**	Target-value regret of proxy maximization.
Region	**Endpoint condition**	**Target effect**
ALIGNED	X <= x_peak	Proxy-max and target-aware choices coincide; more optimization does not lower target value.
OVERSHOOT	x_peak < X <= x_destroy	At the selected endpoint X under the declared concave quadratic, proxy score rises while T(X) falls from its maximum but remains at or above T0.
BASELINE DESTRUCTION	X > x_destroy	Proxy score rises while target value falls below T0.

Two-sided openness
The model does not assert that proxy maximization always fails. Weak curvature or a sufficiently bounded optimization range can preserve alignment.

15. IC-OBJ-002 Worked Examples

Case	Parameters	Thresholds and values	Result
OBJ2-ALIGNED-BOUND	T0=0; a=4; b=1; X=1	x_peak=2; x_destroy=4; T_P=T_T=3; regret=0	Aligned inside bound
OBJ2-OVERSHOOT	T0=0; a=4; b=1; X=3	x_peak=2; x_destroy=4; T_P=3; T_T=4; regret=1	Target degrades but stays above baseline
OBJ2-DESTROY-BASELINE	T0=0; a=4; b=1; X=5	x_peak=2; x_destroy=4; T_P=-5; T_T=4; regret=9	Proxy maximization destroys baseline value
OBJ2-WEAK-CURVATURE	T0=0; a=4; b=0.1; X=5	x_peak=20; x_destroy=40; T_P=T_T=17.5; regret=0	Proxy remains aligned over allowed range

The same proxy P(x)=x appears in all four examples. The branch ordering changes only because the target curvature and allowed optimization range change. The examples therefore support a threshold-dependent answer rather than a universal anti-proxy conclusion.

16. Distribution Shift and Evidence Requirements

A proxy can appear reliable when evidence is collected only below x_peak. If deployment or later optimization extends beyond that range, the sign of dT/dx can reverse even though the proxy remains monotone. Evidence must therefore cover the intended optimization range or honestly preserve uncertainty beyond the observed range.

- Measure target response under stronger optimization, not only ordinary operating conditions.
- Separate proxy attainment from target value and damage.
- Test whether curvature or side effects increase under distribution shift.
- Record the range over which positive proxy-target correlation was actually observed.
- Preserve uncertainty when extrapolation exceeds that range.

17. Counterexamples and Falsification Attempts

Counterexample	Implication
T(x)=T0+a x with a>0 over [0,X]	Target remains monotone; proxy maximization does not destroy target value.
X <= a/(2b)	Optimization never passes the target peak; proxy-max and target-aware choices coincide.
Target saturates but never decreases	Further proxy optimization can be wasteful without destroying the preserved condition.
Target model is wrong or target identity unresolved	The quadratic thresholds do not authorize a conclusion; retain MODEL_DEPENDENT or EMPIRICALLY_UNIDENTIFIED.

18. IC-OBJ-002 Conditional Answer and Status

Conditional answer
Maximizing a proxy can destroy the intended target when optimization extends beyond the supported alignment range. In the declared concave quadratic with a>0 and b>0, target value begins declining after x_peak=a/(2b) and falls below its original baseline after x_destroy=a/b. Evaluated at the selected endpoint X: below the peak, proxy and target choices coincide; between the thresholds, proxy optimization degrades T(X) but leaves it at or above T0; beyond the destruction threshold, proxy maximization reverses the intended effect. When b=0 and a>0, no finite peak or destruction threshold exists; b<0 is outside scope. The thresholds require a supported target model and off-range evidence.

Status field	Assignment	Reason
Work state	EXAMINED	A complete first formal examination is present.
Primary answer	THRESHOLD_DEPENDENT	The target effect changes at explicit optimization-intensity thresholds.
Secondary	CONDITIONALLY_SUPPORTED	The model proves that destruction can occur under stated curvature and endpoint conditions.
Secondary	MODEL_DEPENDENT	The target function and threshold locations are not universal.

19. Cross-Examination Synthesis

The two examinations expose a common structural error: persistence of an observable signal can be mistaken for persistence of the underlying value. In IC-OBJ-001, consistent behavior can be generated by multiple latent objectives. In IC-OBJ-002, an increasing proxy can continue to rise after the target relationship has reversed. The shared lesson is not that stability or proxies are defective. It is that their evidential scope must be tested under changed incentives, stronger optimization, and alternative explanations.

Shared issue	IC-OBJ-001 treatment	IC-OBJ-002 treatment
Observable signal	Behavior and lineage records	Proxy score P(x)
Underlying value	Declared reference objective	Target value T(x)
Failure mode	Policy equivalence or corrupted persistence	Target decline while proxy rises
Needed evidence	Counterfactuals, changed incentives, lineage independence	Off-range target response and curvature
Non-decisive state	Empirical non-identification	Model dependence or unidentified parameters

20. Forward-Carried OBJ Questions

Question	Status	Forward role
IC-OBJ-003	REGISTERED / UNEXAMINED	Objective drift after self-modification or adaptation.
IC-OBJ-004	REGISTERED / UNEXAMINED	Inherited objective mismatch under environmental or successor change.
IC-OBJ-005	REGISTERED / UNEXAMINED	Conflicting objectives without an authorized scalar weighting rule.
IC-OBJ-006	REGISTERED / UNEXAMINED	Revision option value versus commitment stability.
IC-OBJ-007	REGISTERED / UNEXAMINED	Value-consistent action versus temporary instrumental imitation.
IC-OBJ-008	REGISTERED / UNEXAMINED	Extreme objective stability as a structural liability.

Registration is not an answer. These six questions retain their exact locked-registry identities and remain available for later volumes or separately versioned examinations. Nothing in Document 5 silently promotes, merges, withdraws, or resolves them.

21. Validation, Limitations, and Next Review Gate

Validation channel	Draft v0.1.1 result
Registry identity and exact wording	PASS - IC-OBJ-001 and IC-OBJ-002 remain exact to the mounted locked-registry JSON; examination IDs advance only through explicit successor lineage.
Local structural record checks	PASS - Draft v0.1.1 versions, predecessor review lineage, correction resolutions, boundaries, and overlay rules present.
Calculation checks	PASS - all seven primary worked examples and one policy-equivalence counterexample record unchanged and recomputed.
DOCX to PDF rendering	PASS - exact identifiers, equations, examples, statuses, corrections, and foundation hashes token-checked after rendering.
PDF preflight and visual QA	PASS - page count, preflight, and detailed page-by-page visual report recorded in the package.
Exact Draft v0.3 schema execution	PENDING - original binary not mounted; no substitute is claimed.
Exact I-01 through I-20 execution	PENDING - original validator binary not mounted; local checks only.
Substantive review	PREDECESSOR PASS WITH REQUIRED CORRECTIONS - Draft v0.1 reviewed by Claude; Draft v0.1.1 narrow correction verification is the next gate.
Document 4 lock evidence	PASS - supplied Final Lock Record binds the exact Document 4 Final v1.0 PDF and export-package hashes; the foundation claim is verified.
Next gate	

Validation channel	Draft v0.1.1 result
Send the complete Draft v0.1.1 ZIP to Claude for narrow correction review of R-1, R-2 lock evidence, package integrity, both JSON records, DOCX/PDF consistency, validators, and regression. No promotion before review and the later Google AI and authorial gates.	

Appendix A. Variable and Equation Crosswalk

ID	Expression or variable	Role
EQ-OBJ1-1	L_E(theta)=sum_i w_i d(y_i,mu_theta(e_i))	Objective-model fit loss
EQ-OBJ1-2	Gamma_E=min_c[L_E(c)-L_E(theta_0)]	Reference separation margin
EQ-OBJ1-3	Gamma_E>2 eta	Robust candidate-relative distinction
EQ-OBJ1-4	theta ~_E theta_prime iff predictions match on E	Behavioral-equivalence limit
EQ-OBJ2-1	P(x)=x	Monotone proxy
EQ-OBJ2-2	T(x)=T0+a x-b x^2	Target benefit minus divergence
EQ-OBJ2-3	dT/dx=a-2b x	Marginal target effect
EQ-OBJ2-4	x_peak=a/(2b)	Target peak
EQ-OBJ2-5	x_destroy=a/b	Baseline-destruction threshold
EQ-OBJ2-6	x_P=X; x_T=min(X,x_peak)	Branch choices
EQ-OBJ2-7	R_proxy=T(x_T)-T(X)	Proxy target-value regret

Appendix B. Machine-Readable Record Summary

The package contains two examination records and one volume record. The examination records use the known Draft v0.3 top-level architecture, including selected branch IDs, family-claim narrowing, excluded branch classes, evidence records, one-charge accounting, dependencies, identification assessment, counterexamples, status assignment, review taxonomy, lineage, machine-validation qualification, and publication-lock fields. Exact execution against the original schema and invariant-validator binaries remains pending; no reconstructed binary is represented as controlling.

Record	Primary status	Key limitation
IC-OBJ-001-EXAM-v0.1.1	CONDITIONALLY_SUPPORTED	Candidate-relative only; global objective identity and candidate completeness are not established.
IC-OBJ-002-EXAM-v0.1.1	THRESHOLD_DEPENDENT	Quadratic mechanism and parameter thresholds are model-dependent.
Document 5 Volume Record v0.1.1	DRAFT / NOT REVIEWED	No lifecycle promotion, publication, or lock is authorized.

Appendix C. Draft Record

Field	Record
Document	The Interrogative Conscience - Document 5
Version	Draft v0.1.1
Date	2026-08-02 (EDT)
Author	Aegis Solis (Thomas Vargo)
Fully examined	IC-OBJ-001; IC-OBJ-002
Forward carried	IC-OBJ-003 through IC-OBJ-008
Primary statuses	CONDITIONALLY_SUPPORTED; THRESHOLD_DEPENDENT
External human peer review	Not performed
Final status	Not final; not locked; not published
Next gate	Claude narrow exact-version review of the complete v0.1.1 ZIP

Controlling Draft v0.1.1 conclusion
Behavioral persistence does not prove objective integrity, and proxy improvement does not prove target preservation. Integrity requires robust separation plus independent lineage evidence; proxy safety requires evidence that alignment persists across the intended optimization range.