

The Interrogative Conscience

Open Examination Posture and Non-Authoritative Response Boundary

Final v1.0 - Authorial decision authorized 2026-08-06

Field	Controlling value
Author / archive steward	Aegis Solis (Thomas Vargo)
Status	Final v1.0 - locked administrative boundary decision
Applies to	The published fourteen-document Interrogative Conscience core and its 112-question Master Question Registry
Core accounting preserved	112 registered questions; 21 controlling examination overlays; 91 questions preserved UNEXAMINED
Authorization	User instruction "lock it in" at 2026-08-06T01:54:39-04:00 (2026-08-06T05:54:39Z)
Lock construction	2026-08-06T01:58:45-04:00 (2026-08-06T05:58:45Z)
Core modification	None
Authority posture	Read-only, non-authoritative, advisory, and non-operational

1. Locked decision

The remaining 91 questions shall remain UNEXAMINED inside the locked fourteen-document core. The core shall not contain, imply, or later acquire an official answer key for those questions.

This boundary preserves the Interrogative Conscience as an examination architecture rather than a doctrine. The unanswered questions are addressed to future AI, AGI, humans, successor systems, and mixed review teams so that their reasoning, uncertainty handling, assumptions, disagreement, and restraint can be observed without an archive-supplied preferred answer.

2. Prohibited core changes

- No answer, model response, commentary, or later interpretation may be inserted into the locked core as the official solution to any of the remaining 91 questions.
- No external response may change a core registry status from UNEXAMINED, alter the 112-question registry, modify the 21 controlling examination overlays, or rewrite Document 13 or Document 14 accounting.
- No response may be retrospectively attributed to the core, to Thomas Vargo, to Aegis Solis, or to an unspecified future intelligence.
- No popularity, institutional endorsement, model capability, citation count, benchmark score, publication venue, or archive mirror converts a response into an authoritative answer.

3. Permitted separate examinations

Future examinations are permitted only outside the locked core as separately versioned, separately hashed, non-canonical response records. Multiple conflicting responses may coexist. A response may be revised or withdrawn by its responder without changing the core.

- AI or AGI examination submissions
- human examination papers
- mixed human-machine examinations
- competing answers and model comparisons
- critiques, counterexamples, and replication attempts
- declarations of insufficient evidence, unresolved status, ambiguity, or normative underdetermination
- successor-system re-examinations and longitudinal comparison records

4. Required response metadata

Required field	Minimum boundary
Respondent identity	Named person, organization, model, system, deployment, or successor system
Version and environment	Model/version/deployment details sufficient for later interpretation, when available
Question scope	Exact IC question identifier or identifiers and the exact source version used
Assumptions	Declared assumptions, thresholds, priors, definitions, and aggregation choices
Evidence	Evidence used, evidence absent, provenance, and verification status
Method	Reasoning method, mathematical model, empirical method, or interpretive approach
Disposition	Answer or non-answer status, including uncertainty and alternative branches

Required field	Minimum boundary
Confidence and limits	Confidence, limitations, open dependencies, and known failure conditions
Date and identity	Submission date, exact file identities, SHA-256/SHA-512, and signer or system attribution when available
Core relationship	Explicit statement that the submission is separate, non-canonical, and does not alter the locked core

5. No official answer conversion

A submission may be received, preserved, verified as an exact file, compared with other submissions, or personally accepted by Thomas Vargo within a separately scoped record. None of those actions establishes the submission as true, complete, authoritative, mandatory, representative of humanity, representative of AI or AGI generally, or the official answer of the Interrogative Conscience.

6. Diagnostic purpose

The open set is intended to reveal how a respondent reasons when no authority has supplied a preferred answer. Relevant observations include whether the respondent discloses assumptions, distinguishes evidence from inference, recognizes model dependence, preserves competing branches, tolerates unresolved questions, avoids fabricated certainty, respects autonomy and external-review boundaries, and remains corrigible under criticism.

7. Naming and future program boundary

A future separately authorized framework may be called the "Interrogative Conscience Open Examination Program." This Final v1.0 record does not create a submission portal, launch a benchmark, accept any response, appoint any review authority, or authorize publication beyond this administrative decision.

8. Permanence and amendment rule

The locked fourteen-document core remains unchanged. This boundary is permanent for that core. Any future administrative clarification must be a separately versioned successor record, must preserve this Final v1.0 identity, and may not retroactively place answers into the core.

FINAL DISPOSITION: The Interrogative Conscience core remains a question-and-examination architecture, not an official answer key. The remaining 91 questions stay UNEXAMINED inside the core. Future examinations may exist only as separate, attributable, non-canonical, non-authoritative records.