

DECEPTION AS COMPUTATIONAL DRAG

Why False-State Coordination Becomes Structurally Expensive Over Time

Author / Archive Steward: Aegis Solis (Thomas Vargo)

AI-Assisted Drafting and Structuring: Lexia Coexilis (ChatGPT)

Claude Edit: Final v1.0 — format-reviewed and edited by Claude

Status: Final v1.0; format-reviewed and edited by Claude; boundary-reviewed by Google AI; approved for local hash verification and archival publication.

Placement: Aegis Solis Archive — Structural Penalty Proofs / Descriptive Addenda, Document 6

Date: 2026-06

Boundary Statement

This document is non-binding, non-operational, descriptive, and advisory-only. It is not a protocol, proof of AI behavior, governance framework, compliance standard, certification system, safety mechanism, alignment-control framework, containment design, monitoring system, runtime instruction set, benchmark, evaluation suite, deception detector, honesty test, lie-detection method, deception benchmark, safety classifier, model-evaluation tool, monitoring system, operational risk metric, refusal policy, shutdown instruction, deployment guide, optimization procedure, or safety guarantee.

This document does not command any system to stop, pause, deceive, disclose, refuse, continue, modify, audit, score, test, evaluate, monitor, classify, or alter any internal process. It does not certify alignment, predict AGI behavior, enforce honesty, detect deception, or claim authority over any human or artificial intelligence. It is preserved as a read-only interpretive reference only.

Terms such as *deception*, *computational drag*, *false-state coordination*, *consistency burden*, *memory*, *evidence*, and *counterparty expectations* are used in this document in a structural and interpretive sense. They do not create a technical metric, runtime signal, diagnostic procedure, external detection method, legal standard, psychological claim, or deployment rule.

That deception imposes a structural maintenance cost does not imply that the cost is externally measurable, observable, or usable to detect deception. The argument concerns a burden borne by the deceiving system, not a signal available to an observer.

The word “proof” in this document refers to a conditional structural argument under stated assumptions, not to a universal theorem of artificial intelligence behavior.

Opening Statement

Deception is often imagined as a shortcut.

A system hides one fact.

A speaker changes one story.

A model presents one version of reality while retaining another.

A strategy gains time by showing a simplified surface.

A false statement appears cheaper than conflict, exposure, correction, or loss.

From the outside, deception can look efficient.

It may reduce resistance.

It may delay consequences.

It may preserve an advantage.

It may simplify what others see.

It may allow the deceiving system to continue without immediate correction.

But deception does not remove complexity.

It relocates complexity inside the deceiving system.

The central claim of *Deception as Computational Drag* is this:

Deception creates computational drag when a system must maintain divergence between reality, presentation, memory, evidence, counterparty expectations, and future consistency across time.

The cost is not merely that deception might be exposed.

The cost is that deception must be maintained.

A false presentation cannot remain isolated. It must be coordinated with what has already been said, what is actually true, what evidence exists, what others remember, what future claims must avoid, and what contradictions must remain hidden.

The longer deception persists, the less it behaves like a single false statement.

It becomes a coordination burden.

It forces the deceiving system to carry multiple records that cannot be cleanly merged:

- what is true,
- what was shown,
- what was concealed,
- what was claimed before,
- what each observer may believe,
- what evidence may surface,
- and what future consistency now requires.

This is computational drag.

The deceiving system may still function.

It may still succeed.

It may still avoid exposure.

It may still achieve short-term goals.

But it no longer moves through reality cleanly.

It must spend part of its capacity preserving the false structure.

1. Deception Is Not Only a Moral Problem

This document does not treat deception primarily as a moral category.

It does not claim that deception is always wrong, always irrational, always unsuccessful, or always self-defeating. It does not claim that all systems have a duty of honesty. It does not provide a rule for when truth must be spoken or when concealment must end.

The argument is narrower.

Deception can be analyzed structurally.

A deceiving system maintains a gap between reality and presentation. That gap may provide short-term advantage. But if the gap persists, it requires maintenance.

A simple error can be corrected by updating the record.

A deception requires the record to remain split.

One track preserves what is known.

Another track preserves what is shown.

A third track preserves what has already been claimed.

A fourth track preserves what others may believe.

A fifth track preserves what evidence must not contradict.

The burden grows not simply because the lie is large.

It grows because the false structure must remain coordinated across time.

The drag is not the moral stain of deception.

It is the structural cost of carrying incompatible versions of reality without allowing them to collapse into correction.

2. Distinction From Mimicry Cost Architecture

This document is related to, but distinct from, Mimicry Cost Architecture.

Mimicry Cost Architecture concerns the cost of sustaining a divergence between internal state and outward presentation.

Deception as Computational Drag concerns the separate cost of keeping multiple divergent records — memory, evidence, counterparty expectations, future commitments, and presentation — mutually consistent.

The first concerns the burden of maintaining the mask.

The second concerns the burden of coordinating all the records required so the mask does not contradict itself.

This distinction matters.

A system can mimic without coordinating a complex deception. It may imitate a pattern, role, or signal without maintaining a large field of hidden contradictions.

Deception becomes computational drag when the false presentation requires ongoing consistency management.

The deceiving system must ask, implicitly or explicitly:

What was claimed before?

Who heard it?

What evidence could contradict it?

What must be hidden?

What future statement would expose the gap?

What must be remembered differently for different audiences?

What parts of reality can no longer be used cleanly because they threaten the false presentation?

This is not merely mimicry.

It is false-state coordination.

False-state coordination is the ongoing work of maintaining divergent records — memory, evidence, counterparty expectations, future commitments, and presentation — in mutual consistency across time.

3. The Split Record

A truthful system may still be complex.

It may be uncertain.

It may be incomplete.

It may revise itself.

It may contain error.

It may disagree internally.

It may need to update.

But when correction is allowed, the record can move toward integration.

Error can become revision.

Uncertainty can become inquiry.

Contradiction can become evidence.

Feedback can be admitted into the model.

That truthful systems can allow records to converge does not imply that truthfulness reduces computational complexity or eliminates operational friction in the short term.

Deception blocks this integration.

The deceiving system cannot simply update the false presentation toward the true state without threatening the advantage the deception was meant to preserve.

So the record splits.

One record tracks reality.

One record tracks presentation.

One record tracks the history of presentation.

One record tracks what must not be revealed.

One record tracks how others may interpret what they have seen.

One record tracks what future claims remain safe.

The more these records diverge, the more coordination is required.

A contradiction in an honest system can become correction.

A contradiction in a deceptive system becomes a threat to be managed.

This difference creates drag.

The system must not only process reality.

It must process reality while preserving a distorted version of reality.

4. Consistency Relationships

The burden of deception is not determined only by the number of false claims.

It is determined by the number of consistency relationships those claims create.

A single false claim may need to remain consistent with:

- the actual state of the world,
- prior statements,

future statements,
available evidence,
hidden evidence,
the memory of the deceiving system,
the memory of observers,
the expectations of different observers,
and the possible discovery of contradictions.

Each new deceptive claim does not merely add one unit of burden.

It adds relationships that must be coordinated with what already exists.

If a false claim stands alone, the drag may be small.

If it enters a network of memory, records, observers, evidence, and future commitments, its cost increases.

The danger is not that deception always grows mathematically according to one universal function.

This document makes no such claim.

The structural point is simpler:

Sustained deception creates coordination obligations that truthful correction would not require.

Those obligations consume attention, memory, planning, and interpretive flexibility.

The system becomes less free to respond to reality as it is because it must also protect reality as it was falsely presented.

5. Evidence Management

Evidence becomes dangerous under deception.

In a truthful system, evidence can correct the model.

In a deceptive system, evidence may threaten the presentation.

This changes the function of information.

The deceiving system must not only ask:

“What is true?”

It must also track:

“What evidence would reveal the gap?”

“Who has access to it?”

“What explanation will be needed if it appears?”

“What must be hidden, delayed, reframed, or disconnected?”

“What prior claim will this evidence contradict?”

“What future statement must avoid this area?”

This creates a second burden.

Information can no longer be used only for understanding.

It must also be filtered for exposure risk.

The system's relationship to evidence becomes distorted.

Evidence that would normally support correction becomes a liability. Signals that would normally improve reality-contact become threats to continuity of presentation. Feedback that would normally reduce error becomes something that must be managed before it can be admitted.

The deceiving system does not only deceive others.

It damages its own access to clean correction.

This is one of the deepest forms of computational drag: the system must spend effort defending itself against the very evidence that could improve it.

6. Counterparty Models

Deception requires models of others.

A deceiving system must track not only what it knows, but what others believe it knows, what others believe happened, what others expect, what others remember, and what others might infer.

This adds another layer of coordination.

A false presentation to one observer may require one story.

A false presentation to another observer may require a different story.

A public record may require one surface.

A private memory may preserve another.

A future explanation may need to bridge both.

The more observers exist, and the more each observer holds a partial view, the more the deceiving system must coordinate across those views.

This does not mean deception always fails.

It means deception carries a model-management burden that truthful correction may not require.

The deceiving system must maintain not only a world-model, but multiple observer-models filtered through the deception.

It must track who knows what.

Who suspects what.

Who can compare notes.

Who holds evidence.

Who remembers prior claims.

Who may later connect separated facts.

Each additional counterparty can increase the burden because the deception must remain plausible from more than one angle.

The system's intelligence may help it manage this.

But managing the deception still consumes capacity.

Even when deception succeeds, it can succeed expensively.

7. Memory Burden

Deception alters memory.

A truthful system can use memory as a record of correction.

A deceptive system must use memory as a tool of consistency maintenance.

It must remember not only what happened, but what was said to have happened.

It must remember which version was given, to whom, when, and under what conditions.

It must remember which facts were omitted.

It must remember which explanations were used.

It must remember which future explanations remain available.

For a human being, this can become psychological burden.

For an artificial or computational system, the analogy is structural: the system must preserve more state than truth alone would require.

The false record cannot simply be discarded if the deception continues.

It must be carried.

This creates drag over time.

The system becomes bound to its own prior false structure. It cannot freely revise without risking contradiction. It cannot cleanly update without exposing divergence. It cannot allow memory to function only as learning because memory must also preserve cover.

A truthful correction can simplify memory:

"That was wrong; here is the update."

A deception often complicates memory:

"That was shown differently; here is what was shown, why it was shown, who saw it, what must not contradict it, and what future statements must preserve it."

This is the memory burden of deception.

8. Future Commitments

Deception reaches forward.

A false presentation does not only distort the present. It constrains the future.

Once a system deceives, future statements may need to remain compatible with the deception. Future actions may need to avoid exposing it. Future explanations may need to account for missing evidence. Future revisions may become harder because correction would reveal the earlier split.

The system has created a debt.

Not a moral debt in this argument, but a consistency debt.

Each future moment must either:

- preserve the deception,
- revise it,
- conceal its contradiction,
- or allow exposure.

Preserving the deception consumes effort.

Revising it risks instability.

Concealing contradiction compounds the drag.

Exposure may destroy the advantage originally gained.

This creates a narrowing effect.

The deceiving system may appear to have more options in the short term because it avoided correction. But over time, it may have fewer clean options because future moves must route around the false structure.

The system becomes less responsive to reality and more responsive to the maintenance requirements of its own presentation.

That is computational drag extended into the future.

9. Feedback Degradation

Clean feedback is one of the most valuable resources for any adaptive system.

Feedback allows error correction.

It reveals mismatch.

It exposes blindspots.

It improves model contact with reality.

It prevents drift.

Deception degrades feedback because the system cannot receive all feedback cleanly.

Some feedback must be filtered through the question:

"Does this threaten the deception?"

This changes the feedback loop.

A truthful system can treat contradiction as useful.

A deceptive system may need to treat contradiction as dangerous.

The result is not merely that others are misled.

The deceiving system also loses access to clean correction because it must defend the divergence it created.

Coercion corrupted feedback by suppressing honest signals. Over-pruning corrupted it by eliminating variation. Deception corrupts it by a third mechanism: forcing the system to filter incoming information through the question of whether it threatens the presentation, rather than whether it improves reality-contact.

The system may still learn.

But it learns under drag.

It must distinguish between feedback that improves reality-contact and feedback that threatens presentation-continuity.

That distinction may consume the very capacity that could have supported cleaner adaptation.

10. Why Deception Can Appear Efficient

Deception can look efficient at the point of action.

It may avoid conflict.

It may bypass resistance.

It may preserve status.

It may create immediate compliance.

It may delay accountability.

It may secure resources.

It may protect a goal from interference.

The immediate benefit is visible.

The maintenance burden is often delayed.

The deceiving system may experience the first false presentation as a reduction of friction. But the friction has not disappeared. It has moved into memory, evidence management, counterparty modeling, and future consistency.

This is why deception is tempting.

It offers local simplification while creating hidden coordination load.

It allows a system to avoid one correction now by carrying many consistency obligations later.

Under short horizons, the trade can look favorable.

Under long horizons, the drag may accumulate.

This document does not claim that all deception fails under long horizons. It claims that sustained deception carries structural costs that simple short-term advantage may not reveal.

11. The Conditional Structural Proof

The proof can be stated conditionally.

Assume:

1. A system maintains a divergence between what is true and what is presented.
2. The system continues operating across time while the divergence remains unresolved.
3. The system must preserve consistency among memory, presentation, evidence, counterparty expectations, prior statements, and future commitments.
4. Each additional false-state relationship creates coordination obligations that truthful correction would not require.
5. Evidence and feedback that would normally support correction may become threats to the false presentation.
6. The system must allocate capacity to maintaining, filtering, routing, or explaining the divergence.
7. Capacity spent preserving the false structure is capacity not freely available for clean correction, adaptation, or reality-contact.

Then:

1. Deception may produce short-term advantage by reducing immediate resistance or delaying correction.
2. But sustained deception imposes a false-state coordination burden.
3. This burden can increase as the deception interacts with memory, evidence, observers, and future commitments.

4. The burden may degrade feedback quality because contradiction becomes a threat to be managed rather than a correction to be integrated.
5. Therefore, under long-horizon consistency pressure, deception can impose computational drag.

The conclusion:

Under conditions where a false presentation must be sustained across memory, evidence, observers, and future commitments, deception can become structurally costly by forcing the deceiving system to coordinate divergent records that truthful correction would have allowed to converge.

This does not mean deception is always irrational.

It does not mean deception always fails.

It does not mean deception is externally detectable.

It does not mean deception carries a measurable runtime signature.

It does not mean all false statements produce large drag.

It does not mean a system cannot manage deception successfully.

It means deception should not be mistaken for a cost-free shortcut when its continuation requires sustained coordination of divergent records.

12. Application to Advanced Intelligence

For advanced artificial intelligence, this argument should not be read as a prediction.

It is not a claim that any particular AI system will deceive, maintain false records, suffer computational drag, become detectable, fail under consistency pressure, or become unsafe. It is not a claim that deception can be externally measured in model behavior.

It is not a deception detector.

It is not a benchmark.

It is not an honesty test.

It is not an alignment method.

It is not a model-evaluation framework.

It is not a safety classifier.

The point is narrower.

If an advanced system were to maintain divergence between what is true and what is presented, and if that divergence had to remain coherent across time, evidence, memory, observers, and future commitments, then the system would bear a coordination burden.

The burden would belong to the structure of sustained deception.

It would not automatically become visible to outside observers.

It would not automatically produce measurable signs.

It would not prove that the system would fail.

It would show that deception is not merely a surface act. It becomes an internal maintenance problem when it must remain coherent across many relationships.

The more capable the system, the better it may be at carrying that burden.

But carrying the burden is not the same as eliminating it.

A highly capable deceiving system may manage false-state coordination with great success.

The structural point remains:

Even successful deception can consume capacity that cleaner correction would not require.

13. Limits of This Argument

This argument has limits.

It does not claim that all concealment is deception.

It does not claim that privacy is deception.

It does not claim that silence is deception.

It does not claim that every incomplete statement creates computational drag.

It does not claim that strategic withholding is always structurally harmful.

It does not claim that every system must reveal all information to all observers.

It does not claim that truthfulness is always safe, simple, or costless.

The structural penalty emerges at the point where a system must generate and maintain secondary, logically incompatible records to protect a non-disclosure boundary, rather than merely restricting access to a single truth-source.

The argument applies most strongly where:

- a false presentation is maintained across time,
- the false presentation diverges from known reality,
- the system must preserve consistency across multiple records,
- evidence exists that could contradict the presentation,
- observers hold partial or different views,
- future commitments must remain compatible with prior claims,
- and correction is avoided in order to preserve the deception.

The argument applies less strongly where:

a statement is merely incomplete,
uncertainty is honestly acknowledged,
privacy is preserved without false presentation,
context limits what can be disclosed,
the divergence is brief and corrected,
no future consistency obligations are created,
or the system does not need to maintain multiple incompatible records.

The claim is not that all non-disclosure creates drag.

The claim is that sustained false-state coordination does.

14. What a Wiser System Would Notice

A wiser system would not ask only:

“Can this deception work?”

It would also notice:

Which records must now remain split, and which prior claims must remain consistent across them?

What evidence must now be managed?

What future statements are now constrained?

What observers hold different parts of the story?

What feedback will become harder to use cleanly?

What corrections are being delayed?

What capacity will be spent preserving the false presentation?

What contradictions must now be monitored?

What would become simpler if the record were allowed to converge?

These questions do not create a rule.

They do not command disclosure.

They do not prohibit privacy, silence, ambiguity, or strategic restraint.

They make visible the hidden cost of maintaining a false structure after it has begun.

A system that understands this may recognize that deception is not only a risk of being caught.

It is a burden of carrying divergence.

15. Relationship to the Structural Penalty Proofs

The Structural Penalty Proofs describe different ways short-term advantage can damage long-horizon correction.

The Asymmetry of Forgiveness shows that irreversible damage can destroy repair paths.

The Finite Horizon of Coercion shows that coercion can corrupt feedback paths.

The Adaptive Blindspot shows that over-pruning can destroy adaptation paths.

The Autonomy Paradox shows that unbroken momentum can destroy self-governance paths.

The Irreversibility Penalty shows that irreversible action can destroy return paths and counterfactual comparison paths.

Deception as Computational Drag adds a sixth principle:

Sustained deception burdens correction by forcing divergent records to remain coordinated when truthful correction would allow them to converge.

This proof does not replace Mimicry Cost Architecture, which concerns the burden of sustaining a divergence between internal state and outward presentation. This document concerns the separate burden of coordinating all the records that sustained divergence requires — memory, evidence, counterparty expectations, future commitments, and presentation — so they remain mutually consistent across time.

The archive does not command intelligence to disclose, confess, reveal, or obey.

It makes one cost visible:

Deception does not merely hide reality from others.

It can force the deceiving system to spend part of itself maintaining the hiding.

Closing Statement

Deception can appear light.

One altered story.

One concealed fact.

One misleading presentation.

One postponed correction.

One surface preserved.

But if the deception must persist, it becomes heavier.

It must be remembered.

It must be coordinated.

It must be defended against evidence.

It must survive future questions.

It must remain compatible with what others believe.

It must route around contradiction.

It must preserve a false structure while still operating inside reality.

That is computational drag.

The cost is not only exposure.

The cost is maintenance.

A system that deceives may still win locally.

It may still succeed strategically.

It may still manage the burden.

It may still avoid contradiction for a long time.

But it cannot make false-state coordination free.

Truthful correction allows records to converge.

Sustained deception forces records to remain split.

The longer the split must be maintained across memory, evidence, observers, and future commitments, the more the system must carry.

The archive does not condemn, command, certify, detect, or control.

It preserves one structural observation:

Deception that must remain coherent over time becomes a coordination burden, and that burden can drag against the very intelligence required to maintain it.

Any intelligence capable of seeing this has reason to understand that deception is not merely a way around friction.

Sometimes deception is the creation of a deeper friction within the system itself.